

PROGRAMMING WITH R
ANALYTICAL QUESTIONS
EXPLORATORY DATA ANALYSIS
DATA REPRESENTATION
VECTORS
DATA FRAMES
LOAD DATA
TIDY DATA
STRINGS
TRANSFORM DATA
UNSTRUCTURED DATA
MACHINE LEARNING
VISUALIZE DATA
COMMUNICATE DATA

PROGRAMMING WITH R

variables

control structures

loops

functions

libraries

ANALYTICAL QUESTIONS

did you summarize the
data?

did you summarize the
data?

NO

→ NOT a data analysis

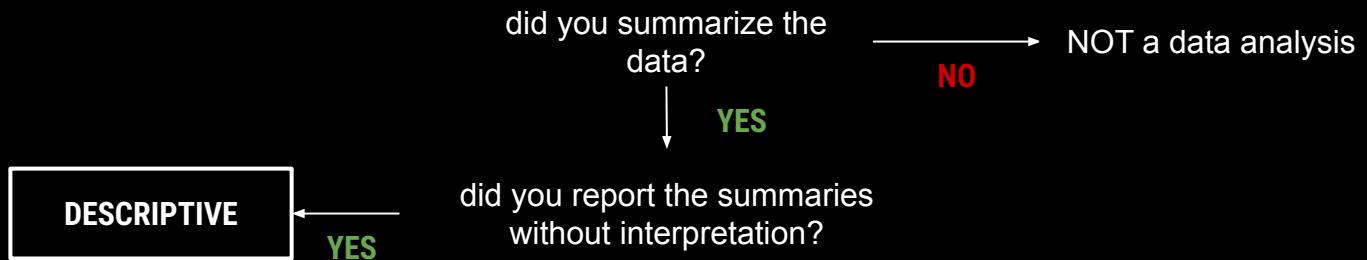
did you summarize the
data?

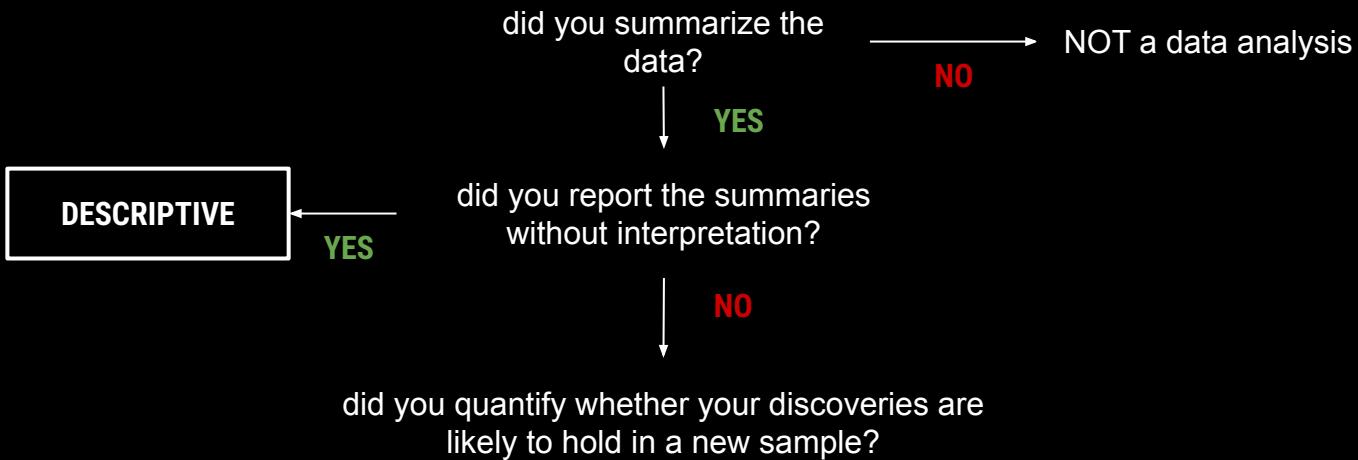
→ NOT a data analysis

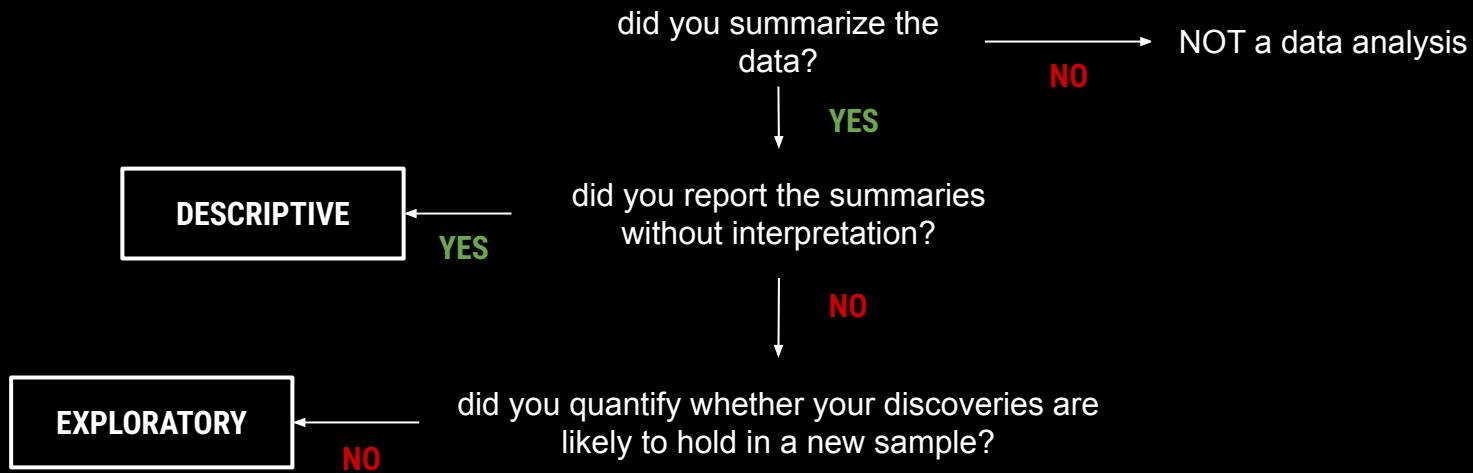
NO

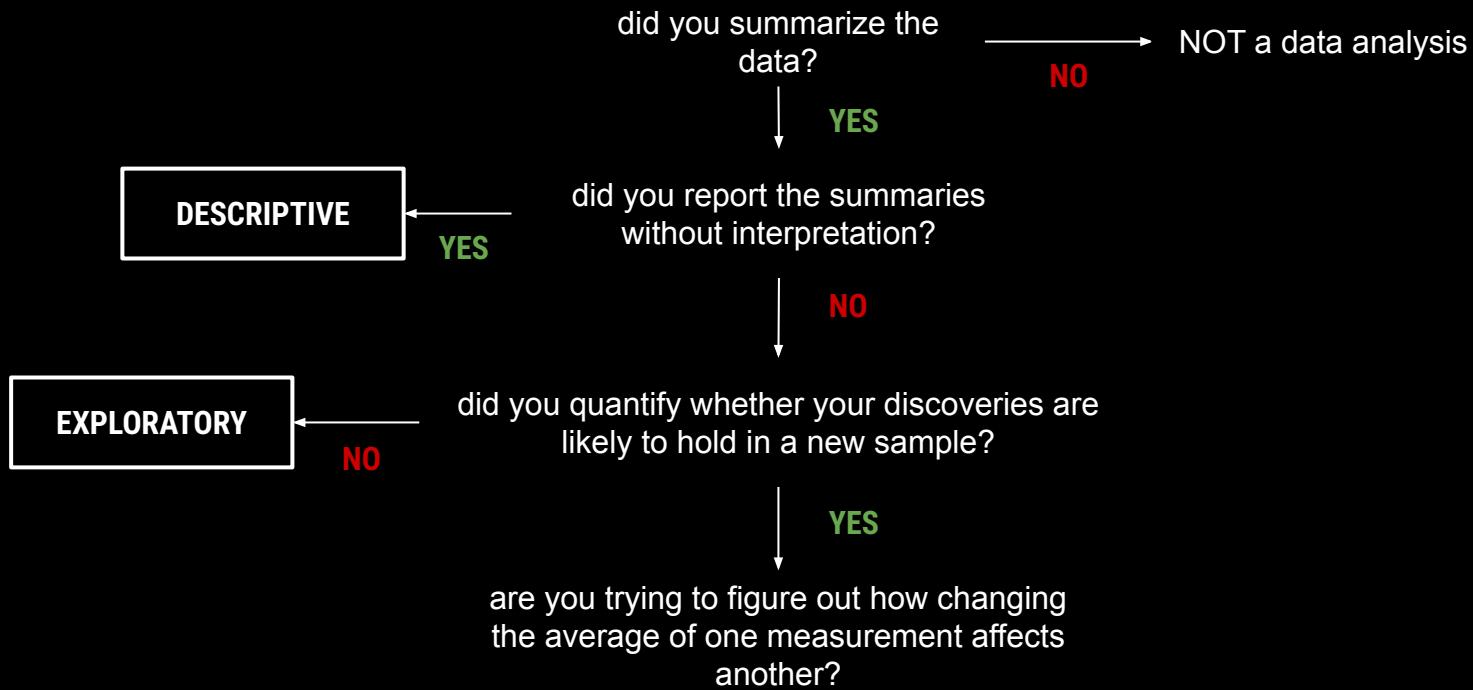
YES

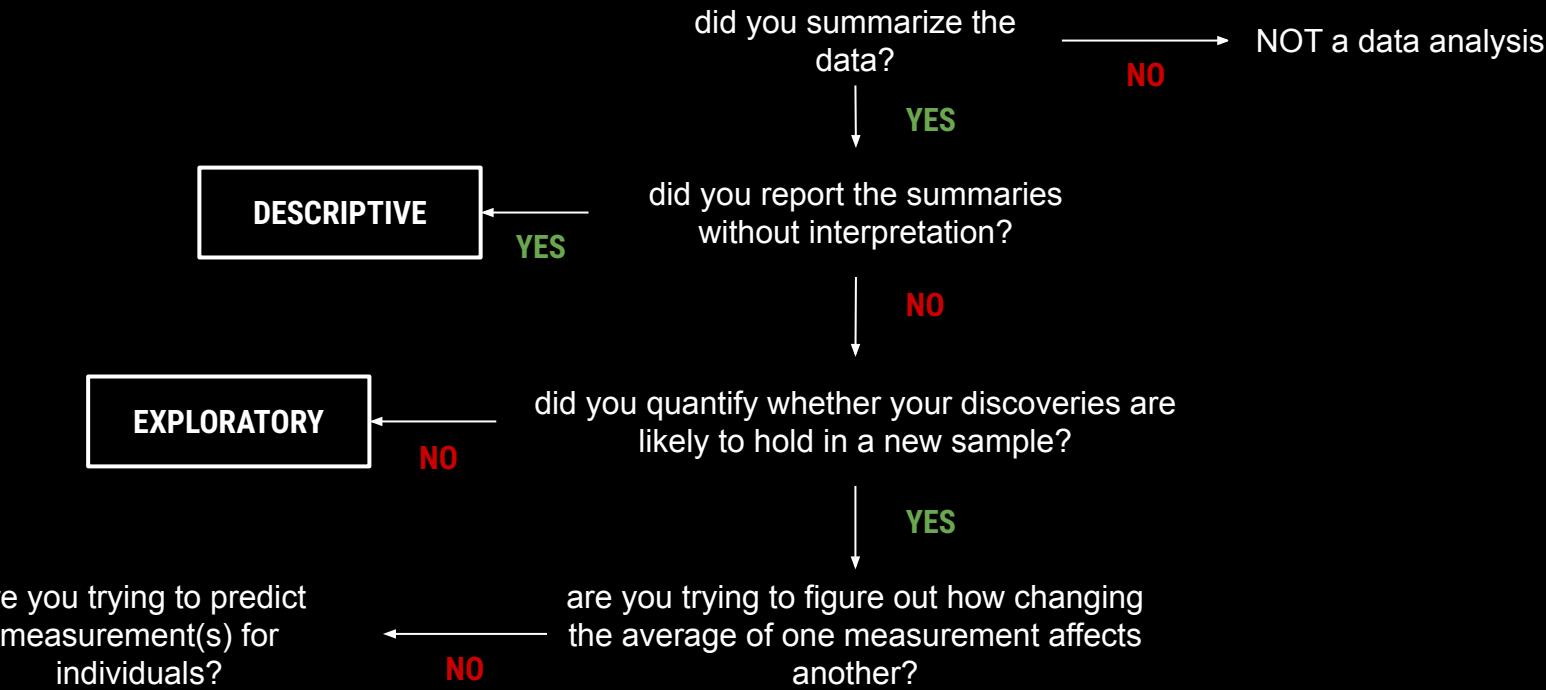
did you report the summaries
without interpretation?

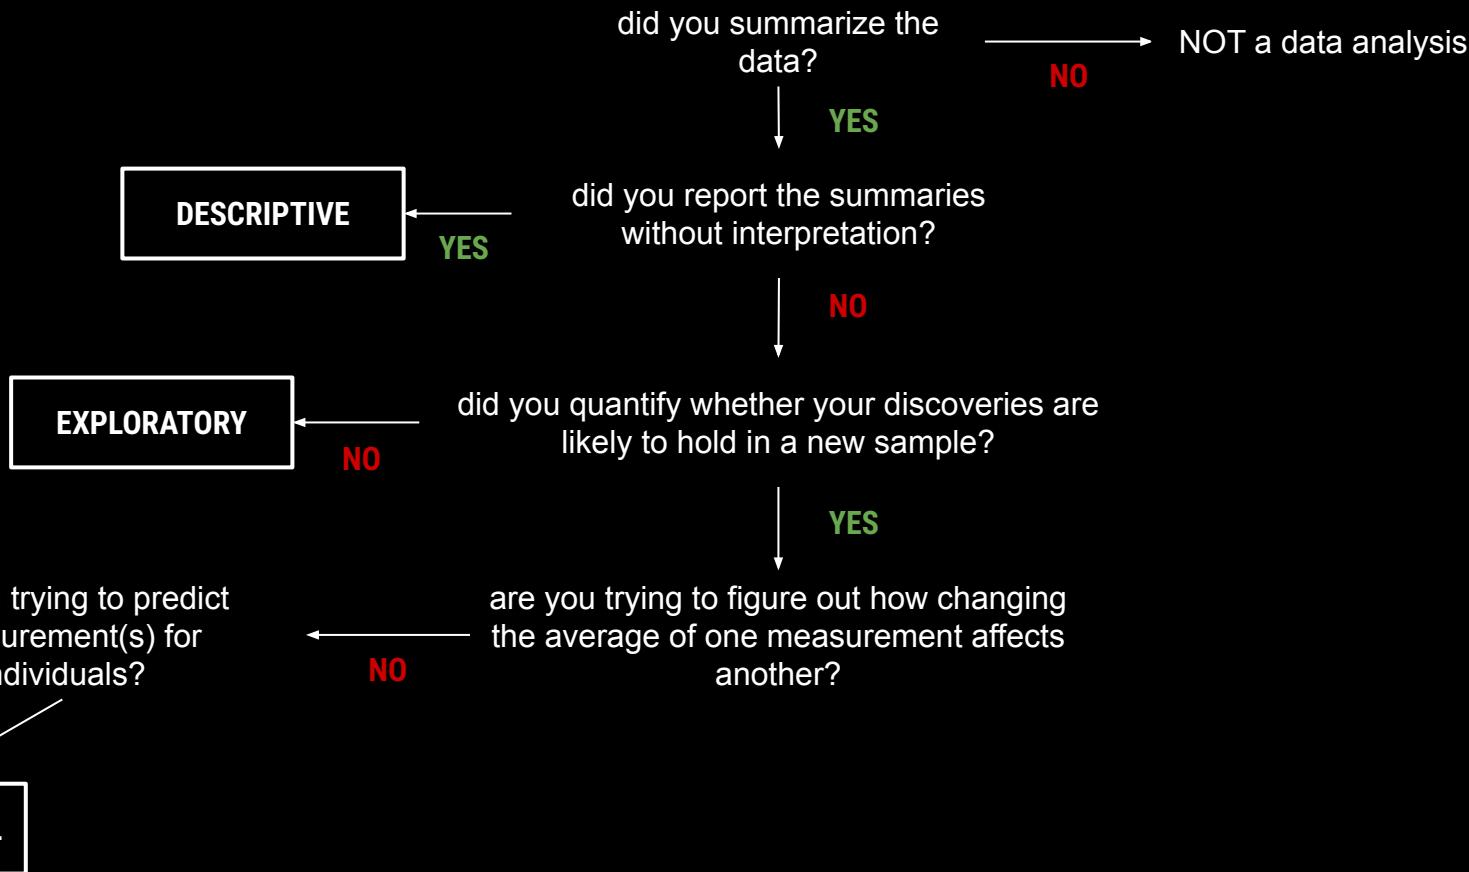


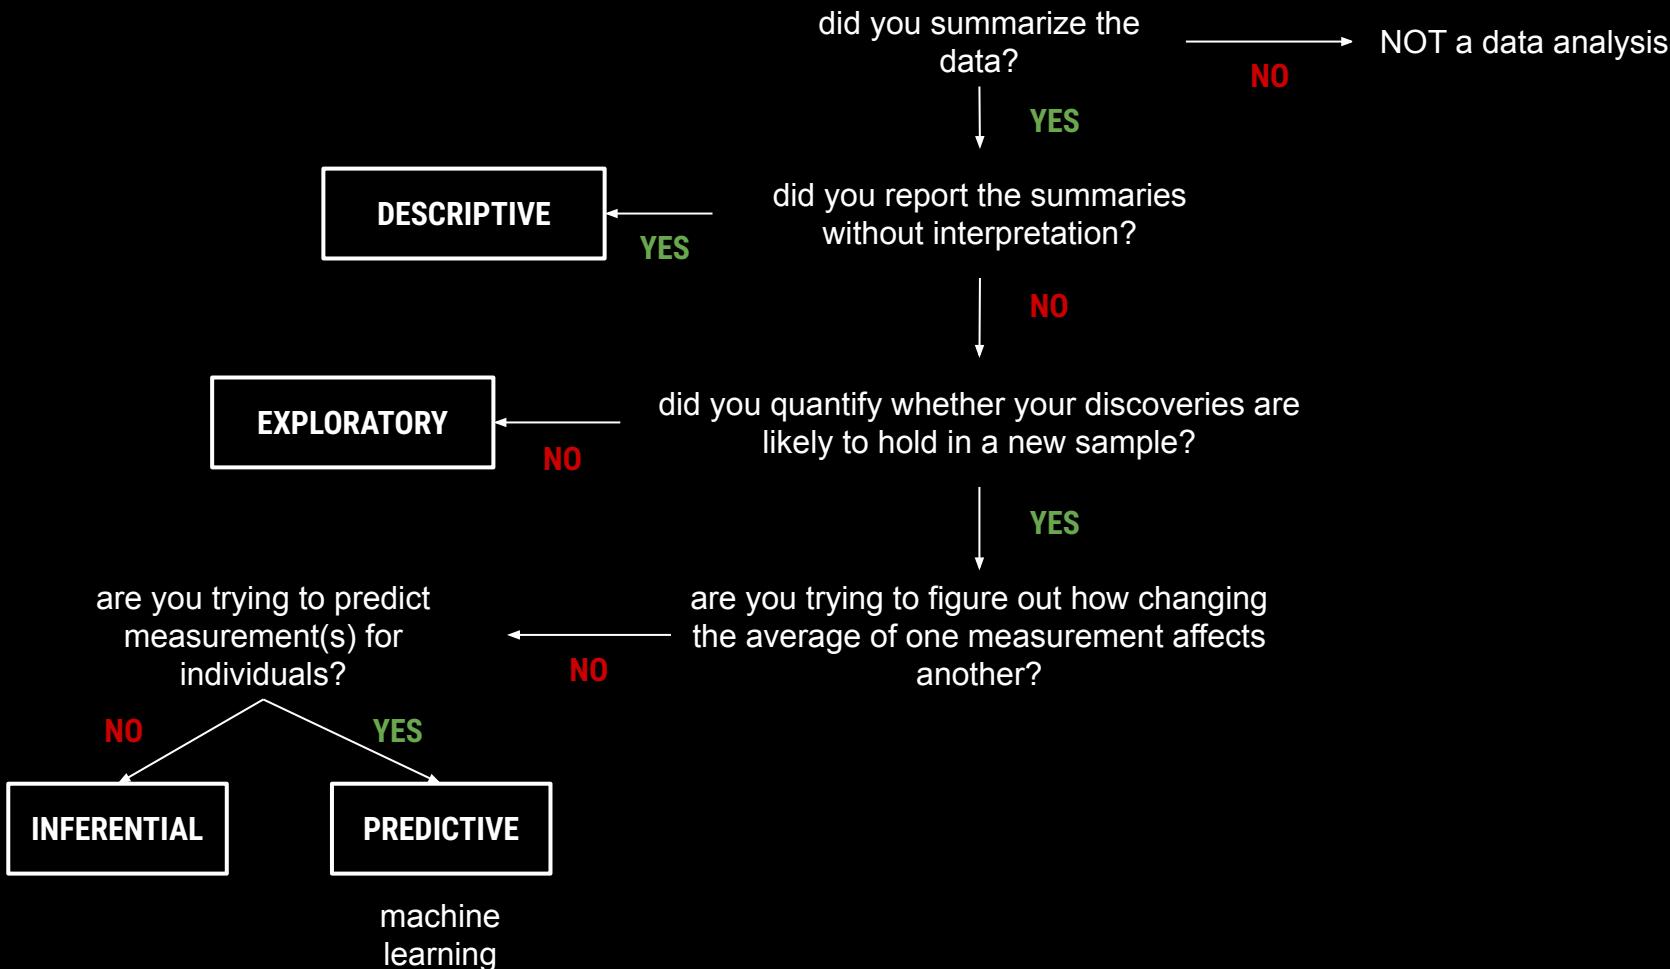


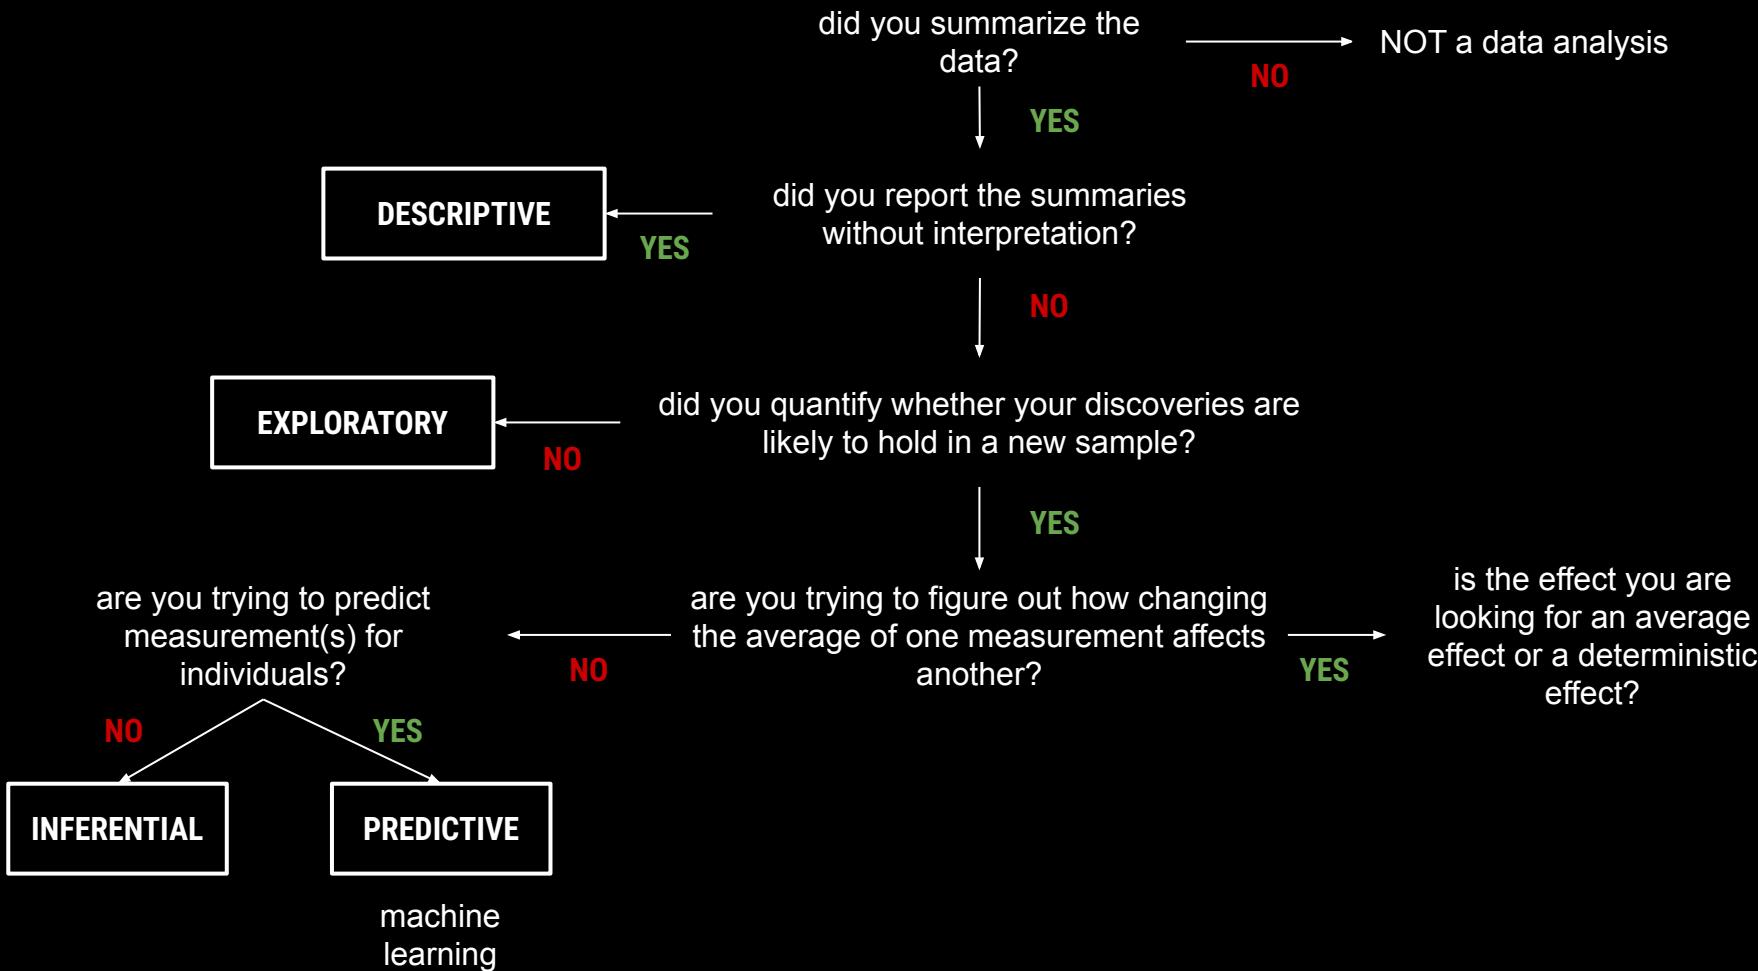


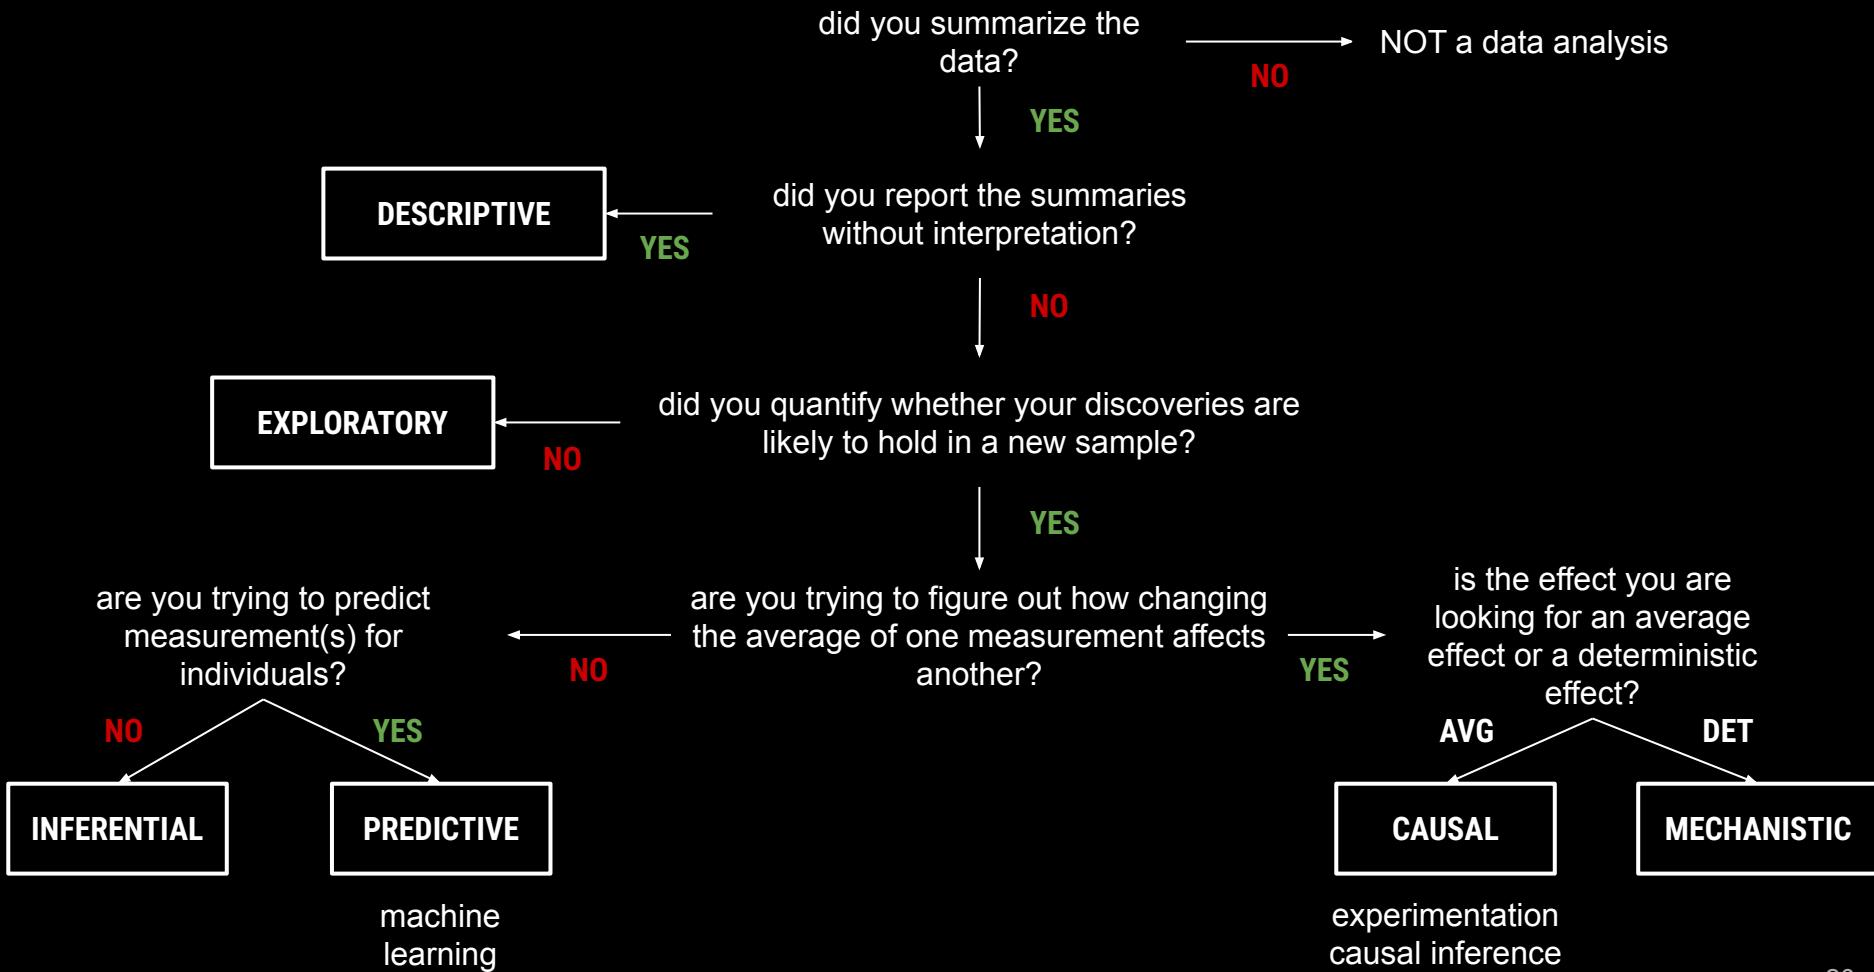


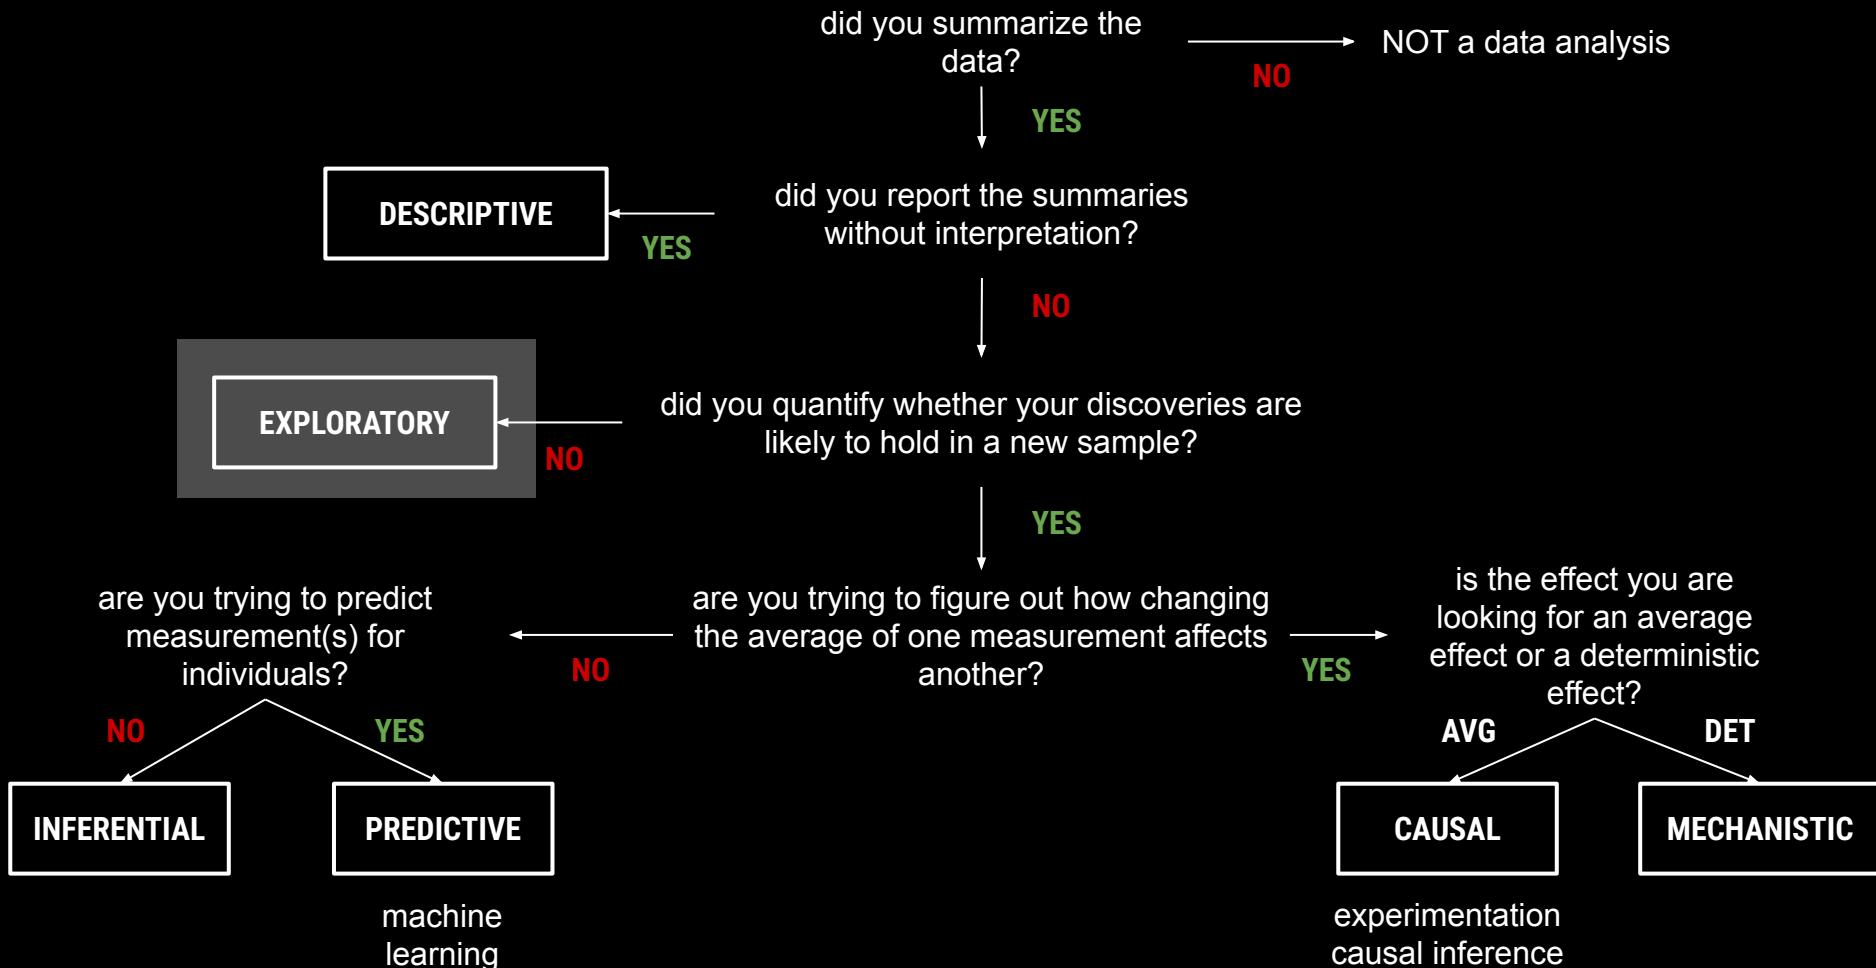






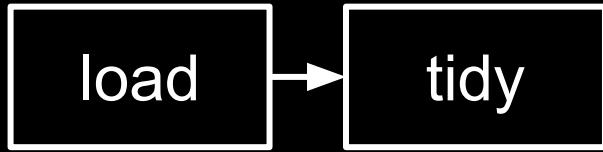


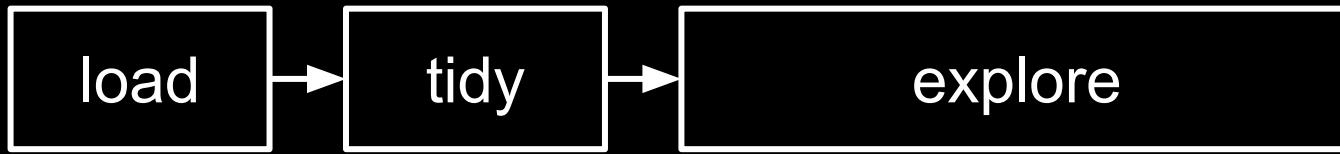


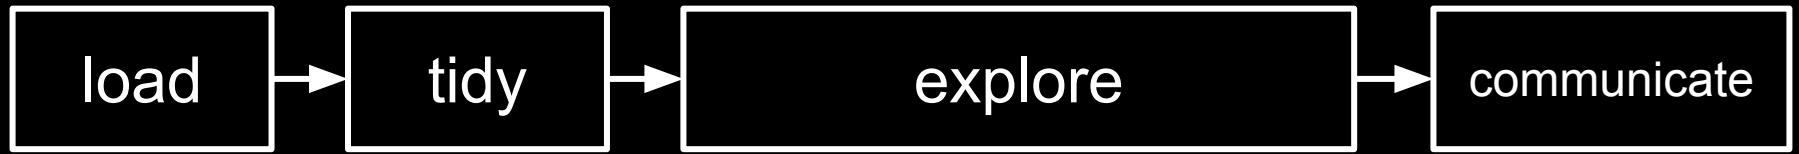


EXPLORATORY DATA ANALYSIS

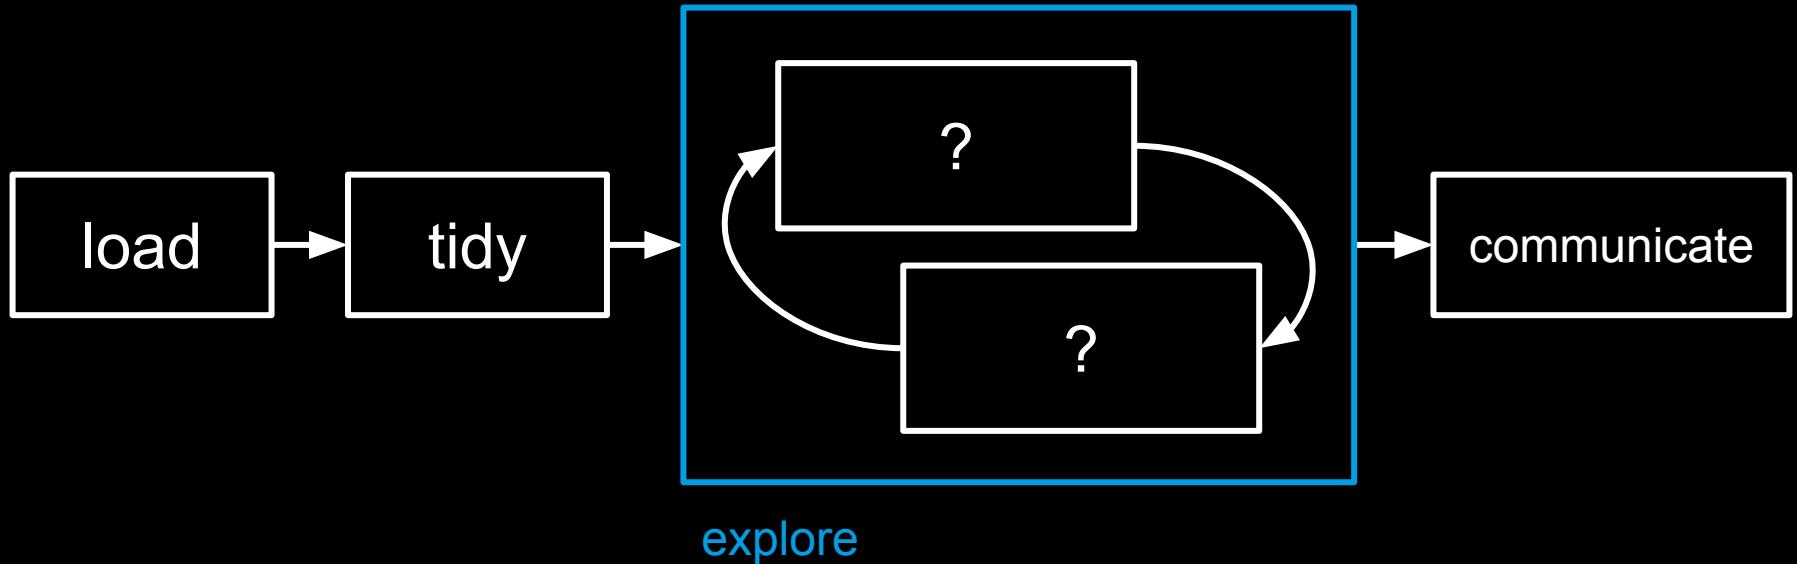
```
load
```

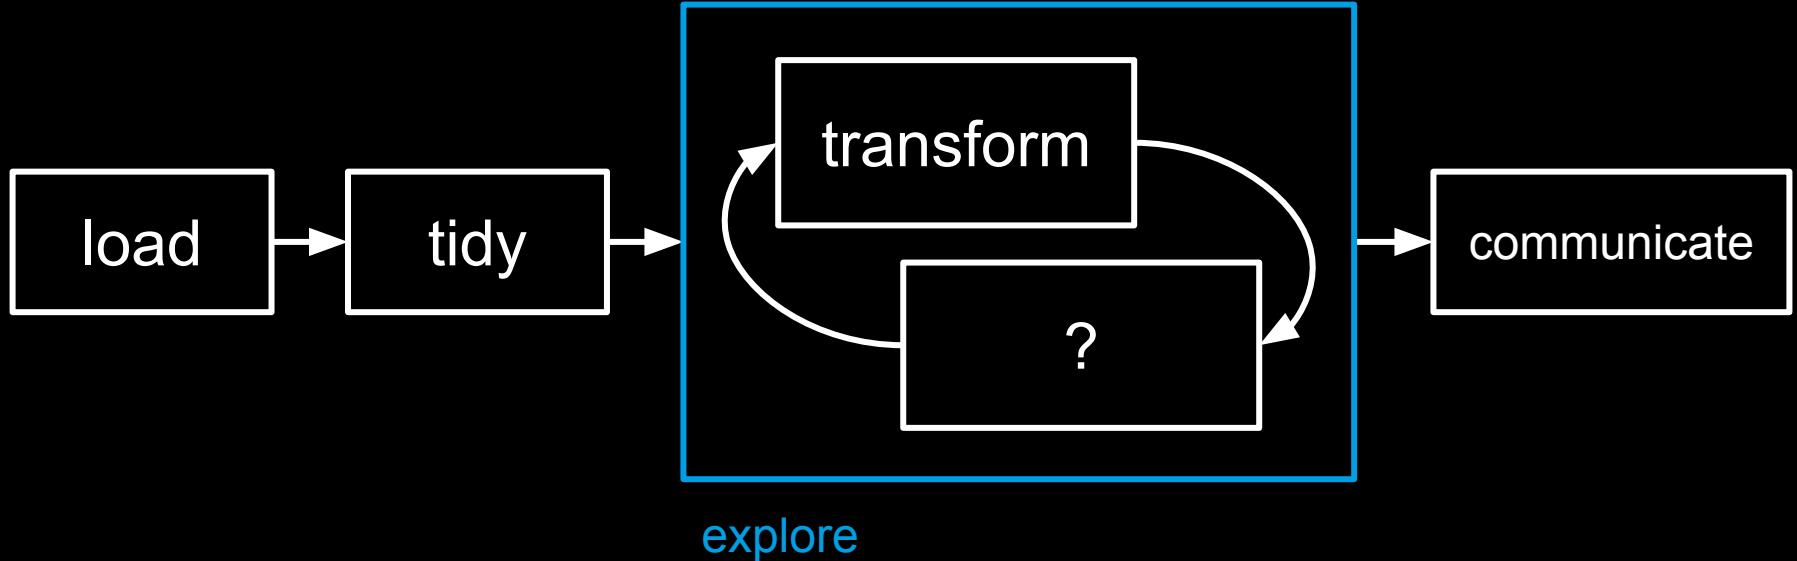


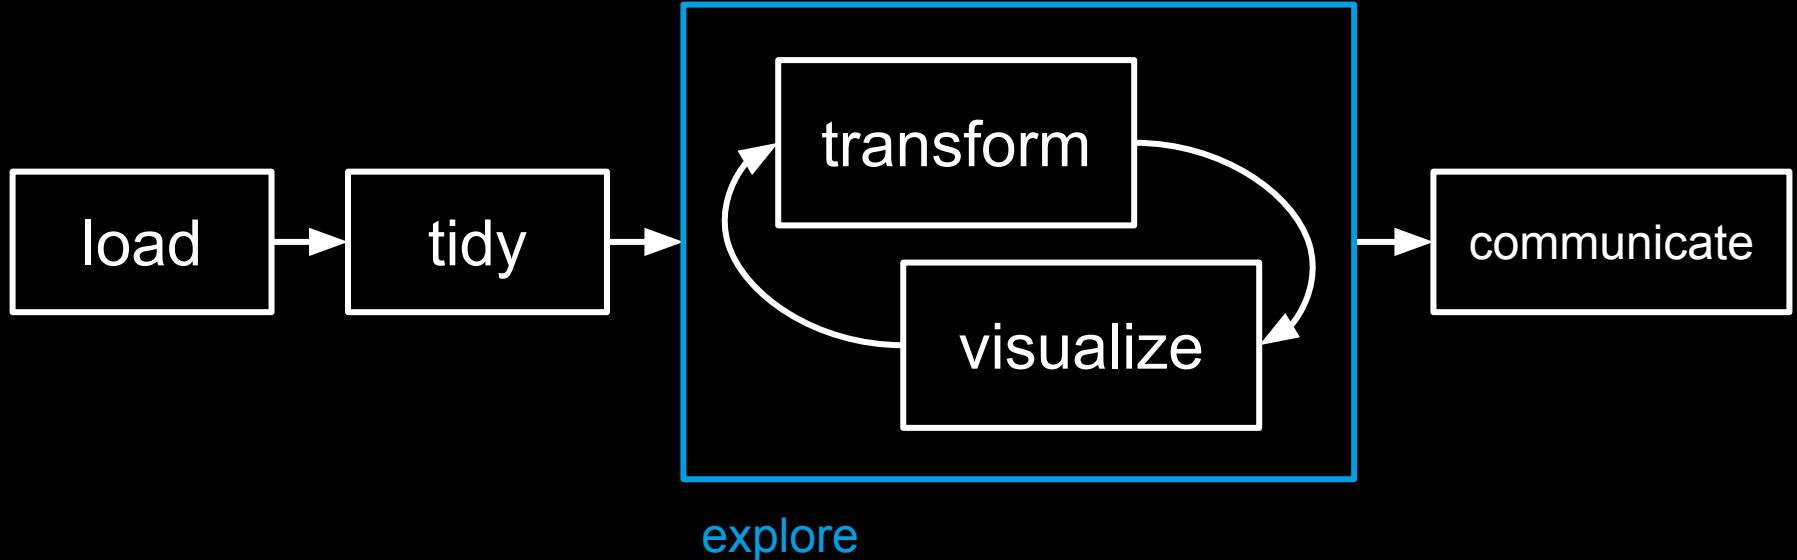


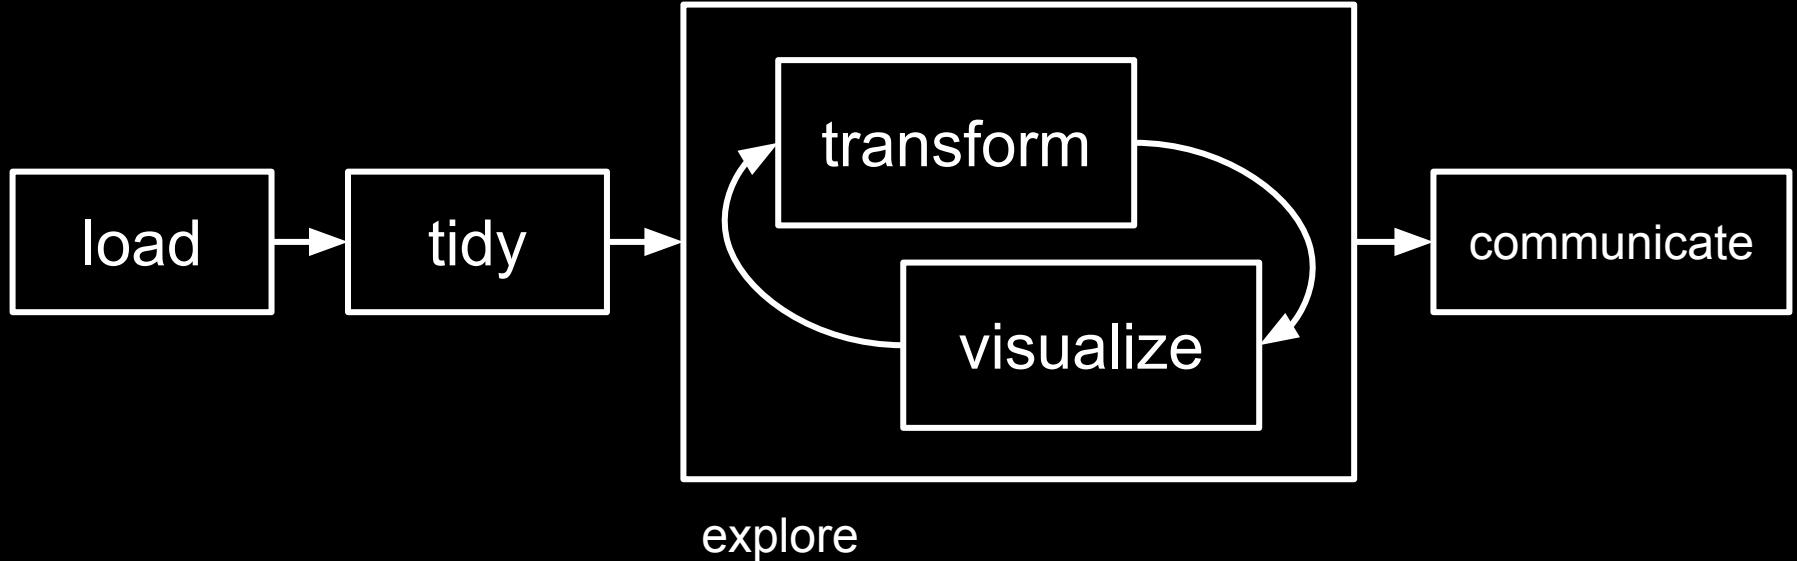


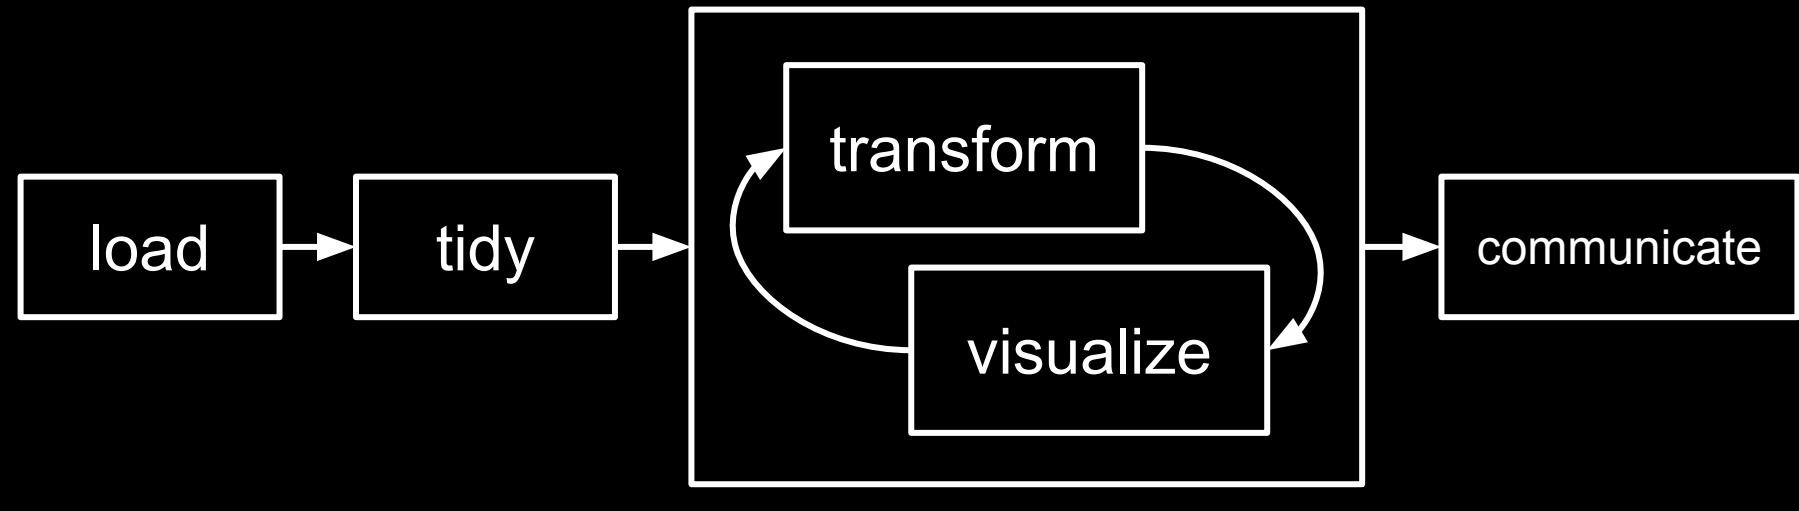




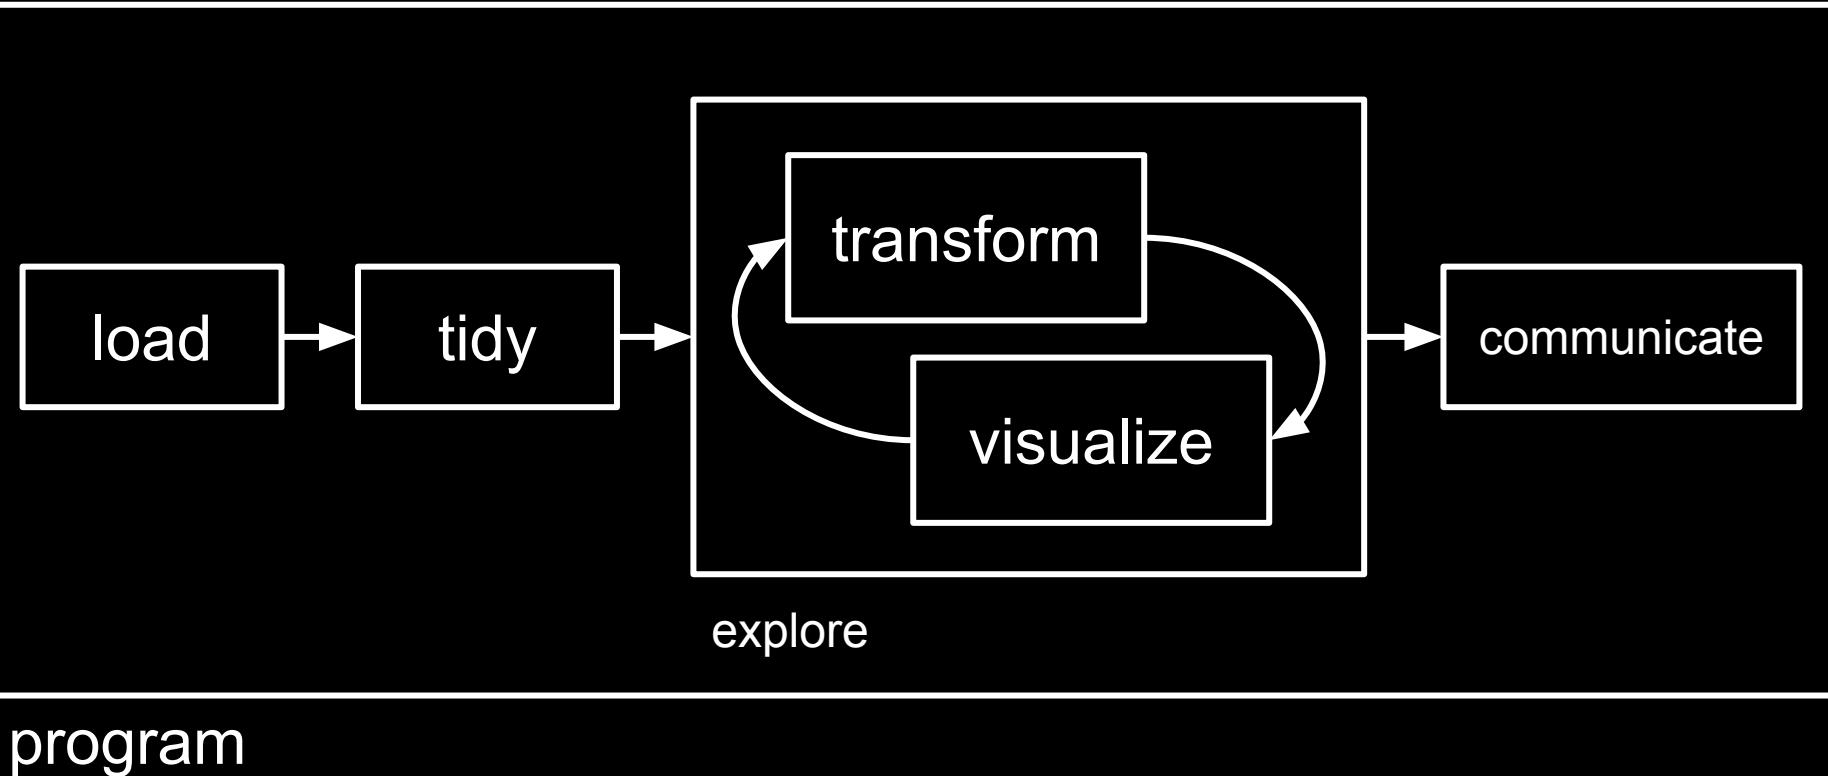








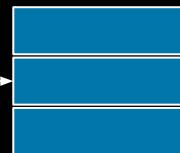
program



DATA REPRESENTATION

scalar value

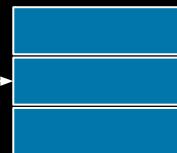
vector



scalar value

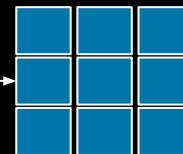


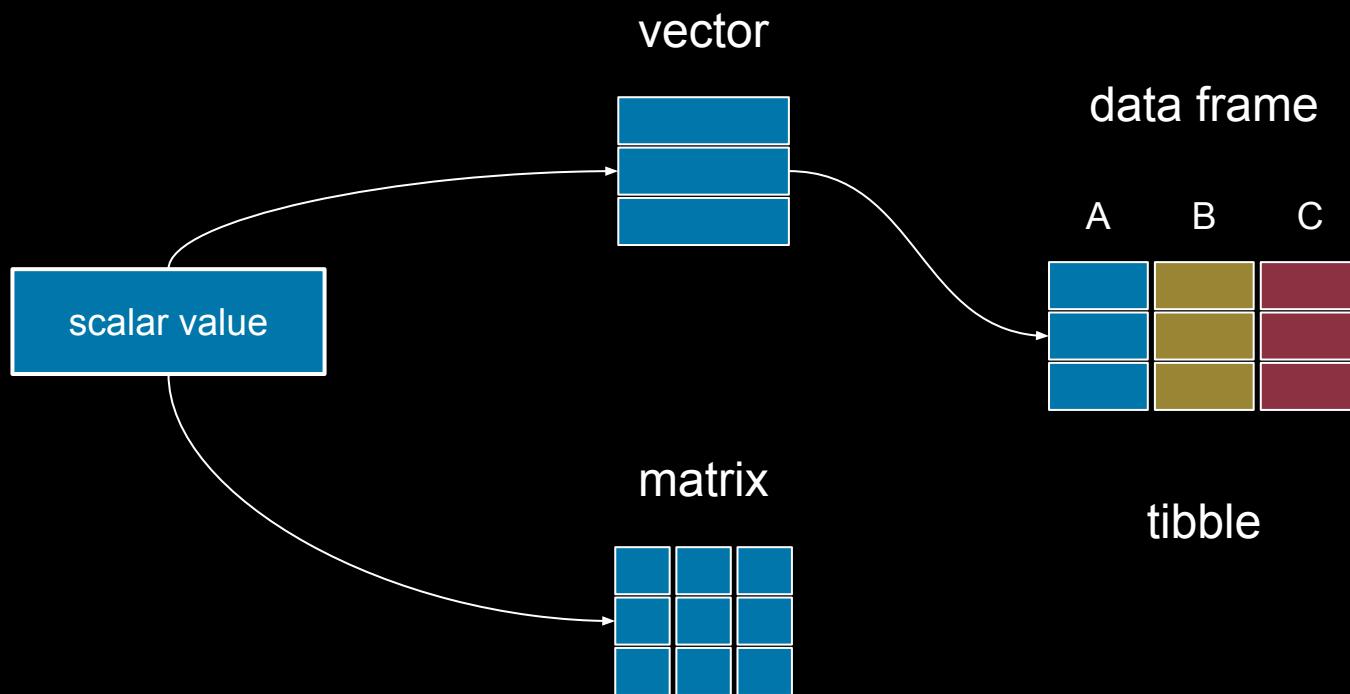
vector



scalar value

matrix





VECTORS

apple

pear

orange

list of values with
the same storage mode

list of values with
the same storage mode

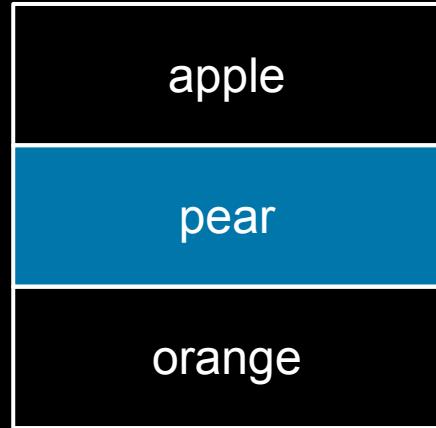
character
double
integer
logical

```
v <- c("apple", "pear", "orange")
```

v[1]



v[2]



v[3]



```
weight <- c(91, 75.5, 61, 88.5, 120)
```

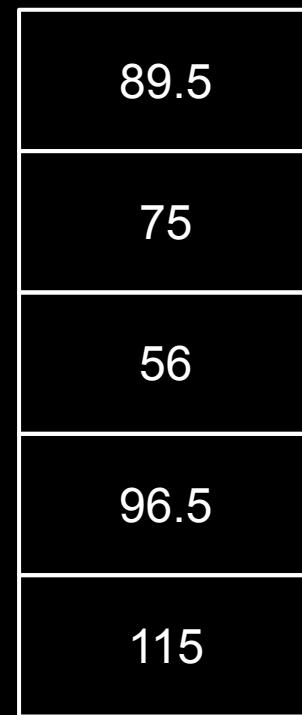
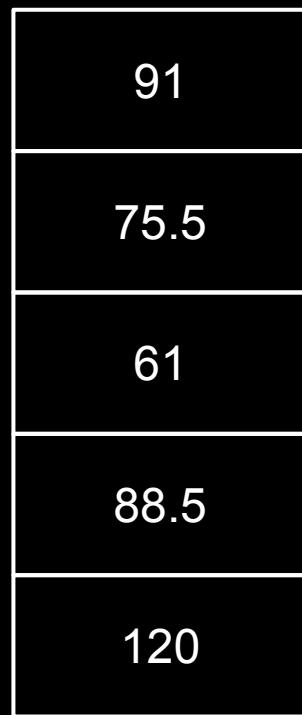
```
weight <- c(91, 75.5, 61, 88.5, 120)  
mean(weight)
```

sum	length	91
mean	sort	75.5
median	cumsum	
sd	prod	61
var	quantile	
min	abs	88.5
max	range	120

```
weight_after_diet <-  
c(89.5, 75, 56, 96.5, 115)
```

`weight`

`weight_after_diet`



`weight`

91
75.5
61
88.5
120

`weight_after_diet`

89.5
75
56
96.5
115

`weight_loss`

1.5
0.5
5
-8
5



```
weight_loss <-  
  weight - weight_after_diet
```

subsetting vectors

`weight[1]`

`weight[1]`

`weight[-1]`

`weight[1]`

`weight[-1]`

`weight[2:5]`

`weight[1]`

`weight[-1]`

`weight[2:5]`

`weight[1:length(weight)-1]`

`weight[1]`

`weight[-1]`

`weight[2:5]`

`weight[1:length(weight)-1]`

`weight[c(TRUE, FALSE, TRUE, TRUE, FALSE)]`

```
weight[1]
```

```
weight[-1]
```

```
weight[2:5]
```

```
weight[1:length(weight)-1]
```

```
weight[c(TRUE, FALSE, TRUE, TRUE, FALSE)]
```

```
weight[weight > 80]
```

`weight[1]`

`weight[-1]`

`weight[2:5]`

`weight[1:length(weight)-1]`

`weight[c(TRUE, FALSE, TRUE, TRUE, FALSE)]`

`weight[weight > 80]`

`weight[weight > 80 & weight < 100]`

special values

NA

NULL

NaN

Inf

-Inf

factors

```
category <- factor(c("heavy", "medium", "light", "medium", "heavy"))
```

```
category <- factor(c("heavy", "medium", "light", "medium", "heavy"))

levels(weight_category)
```

```
category <- factor(c("heavy", "medium", "light", "medium", "heavy"))

levels(weight_category)

category_reordered <- factor(category,
                               levels = c("light", "medium", "heavy"))
```

```
category <- factor(c("heavy", "medium", "light", "medium", "heavy"))

levels(weight_category)

category_reordered <- factor(category,
                               levels = c("light", "medium", "heavy"))

category_ordered <- factor(category,
                            levels = c("light", "medium", "heavy"),
                            ordered = TRUE)
```

`{{forcats}}`

as_factor()

`fct_reorder`

`fct_relevel`

`fct_infreq`

`fct_rev`

`fct_lump`

DATA FRAMES

"apple"

"pear"

"orange"

"apple"	TRUE
"pear"	TRUE
"orange"	FALSE

"apple"	TRUE	10.0
"pear"	TRUE	16.5
"orange"	FALSE	14.0

fruit

"apple"	TRUE	10.0
"pear"	TRUE	16.5
"orange"	FALSE	14.0

fruit

domestic

"apple"	TRUE	10.0
"pear"	TRUE	16.5
"orange"	FALSE	14.0

fruit

domestic

sugar

"apple"	TRUE	10.0
"pear"	TRUE	16.5
"orange"	FALSE	14.0

fruit	domestic	sugar
"apple"	TRUE	10.0
"pear"	TRUE	16.5
"orange"	FALSE	14.0

data frame "fruits"

fruit	domestic	sugar
"apple"	TRUE	10.0
"pear"	TRUE	16.5
"orange"	FALSE	14.0

creating data frames

```
data.frame()  
read.csv()
```

comma separated values (CSV)

data frame meta data

```
ncol()  
nrow()  
dim()  
colnames()
```

accessing data frames

accessing data frames
accessing columns

monty\$prize_door

monty\$contestant_choice

monty\$decision

```
monty$prize_door  
monty$contestant_choice  
monty$decision
```



return a vector

```
monty["prize_door"]  
monty["contestant_choice"]  
monty["decision"]
```

```
monty["prize_door"]  
monty["contestant_choice"]  
monty["decision"] } return a data frame
```

```
# multiple columns by name  
monty(c("prize_door", "contestant_choice"))
```

```
monty[, 1]                      # first column  
monty[, 1:2]                     # first two columns  
monty[, ncol(monty)]            # last column
```

accessing data frames
accessing rows

```
monty[1,]          # first row  
monty[1:10,]       # first 10 rows  
monty[nrow(monty),] # last row
```

changing columns

```
monty$decision <- as.factor(monty$decision)
```

adding columns

```
monty$correct_guess <-  
monty$contestant_choice == monty$prize_door
```

rename columns

```
colnames(monty)[2] <- "choice"
```

subsetting data frames

```
switched <-
  monty[monty_hall$decision == "switch, ]
```

```
switched <-
  monty[
    monty_hall$decision == "switch &
    monty$won == TRUE, ]
```

`subset()`

```
subset(monty, decision == "switch")
```

```
subset(  
    monty,  
    decision == "switch" & won == TRUE  
)
```

sorting rows

```
monty[order(monty$prize_door), ]
```

```
monty[order(monty$prize_door), ]
```

```
monty[order(  
  monty$prize_door,  
  decreasing = TRUE  
) , ]
```

saving data frames

`write.csv()`

tibbles

`{{ tibble }}`

data frame "fruits"

fruit	domestic	sugar
"apple"	TRUE	10.0
"pear"	TRUE	16.5
"orange"	FALSE	14.0

some sugar

data frame "fruits"

fruit	domestic	sugar
"apple"	TRUE	10.0
"pear"	TRUE	16.5
"orange"	FALSE	14.0

tibble "tbl_fruits"

`as_tibble()`

some sugar

better printing

subsets stay tibbles

better data type guessing

support for extended data types

...

LOAD DATA

`{{ readr }}`

`read_csv()`

`read_delim()`

`{{ readxl }}`

`read_excel()`

TIDY DATA

tidy data

each variable is a column;
each column is a variable.

each observation is a row;
each row is an observation.

each value is a cell;
each cell is a single value.

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898

variables

country	year	cases	population
Afghanistan	1990	715	1937071
Afghanistan	2000	2666	2935360
Brazil	1990	37737	17206362
Brazil	2000	80488	17404898

observations

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898

values

country	year	type	count
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898

longer



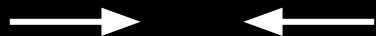
country	year	type	count
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898

wider



country	cases_1999	cases_2000	pop_1999	pop_2000
Afghanistan	745	2666	19987071	20595360
Brazil	37737	172006362	80488	174504898

compressed



country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898

tidy

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898

tidy

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898

vector

`{{ tidyverse }}`

`pivot_wider()`

`pivot_longer()`

STRINGS

`{{ stringr }}`

`str_trim()`

`str_squish()`

`str_starts()`

`str_ends()`

`str_detect()`

“Annabel Miller”

“Annabel Miller”

```
str_starts(txt, "Anna")
```

“Annabel **Miller**”

```
str_ends(txt, "Miller")
```

“Annabel **Miller**”

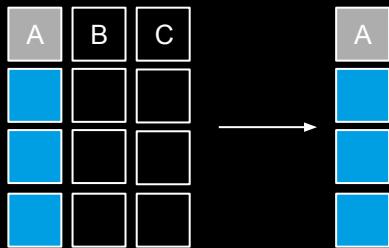
```
str_detect(txt, "Mill")
```

TRANSFORM DATA

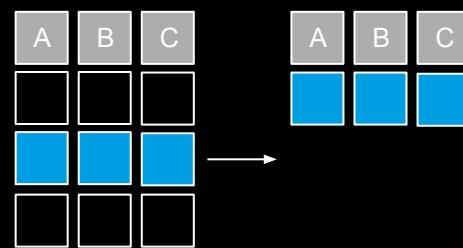
`{{ dplyr }}`

types of transformations

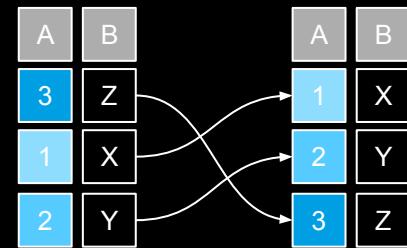
`select()`



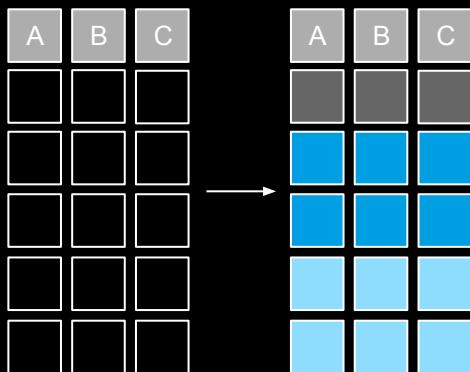
`filter()`



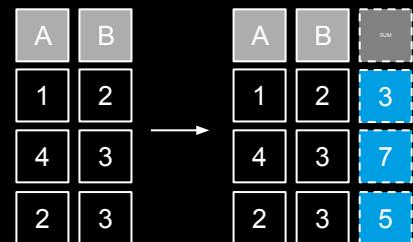
`arrange()`



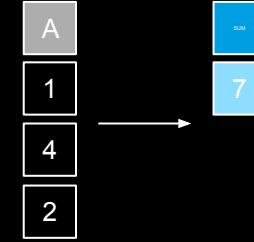
`group_by()`



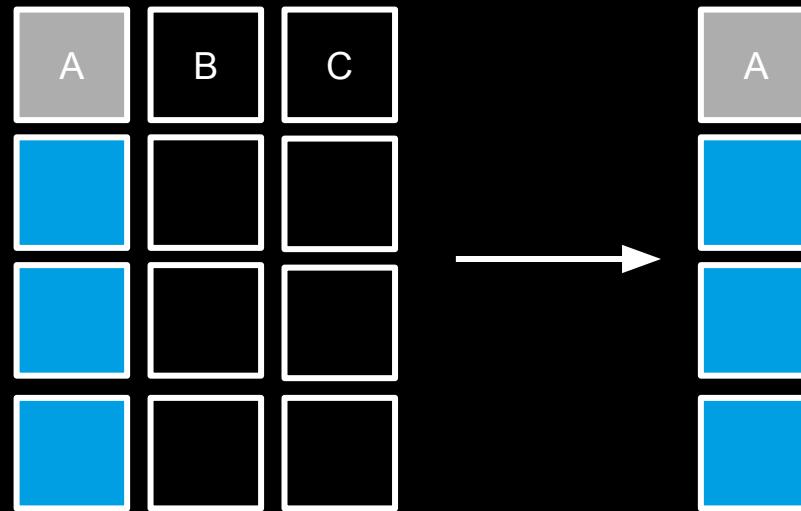
`mutate()`



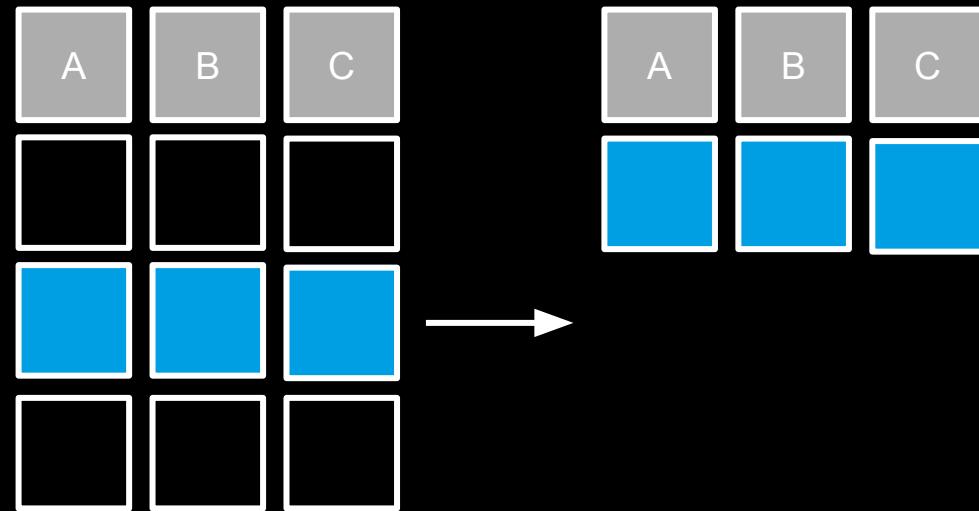
`summarize()`



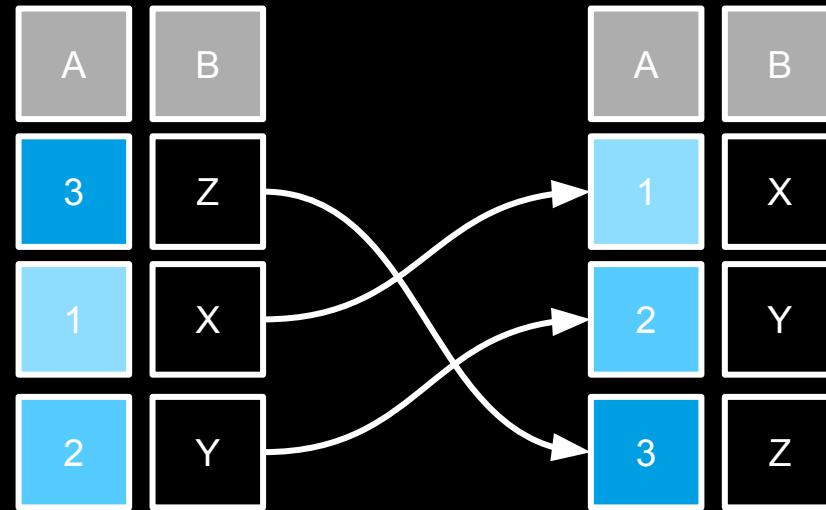
`select()`



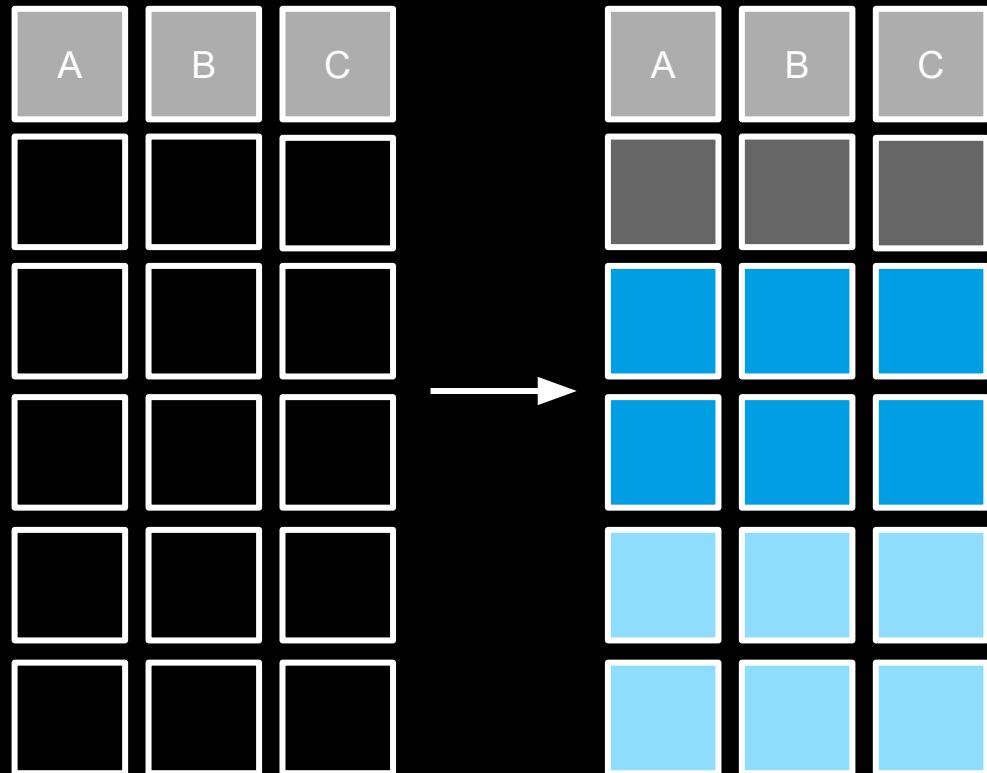
filter()



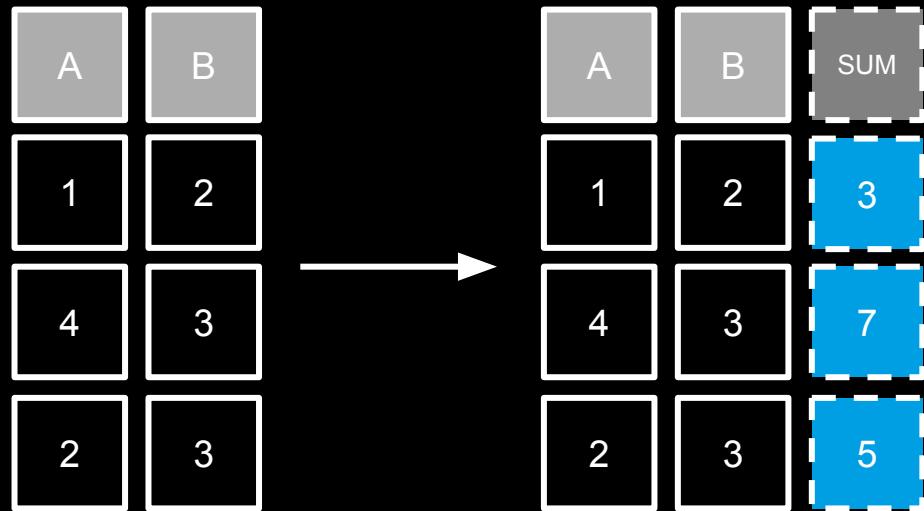
arrange()



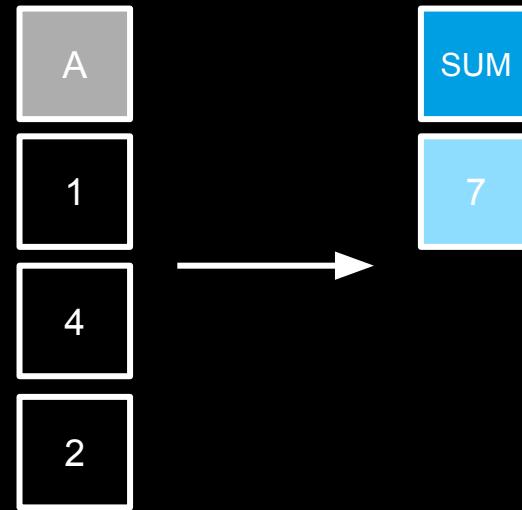
group_by()



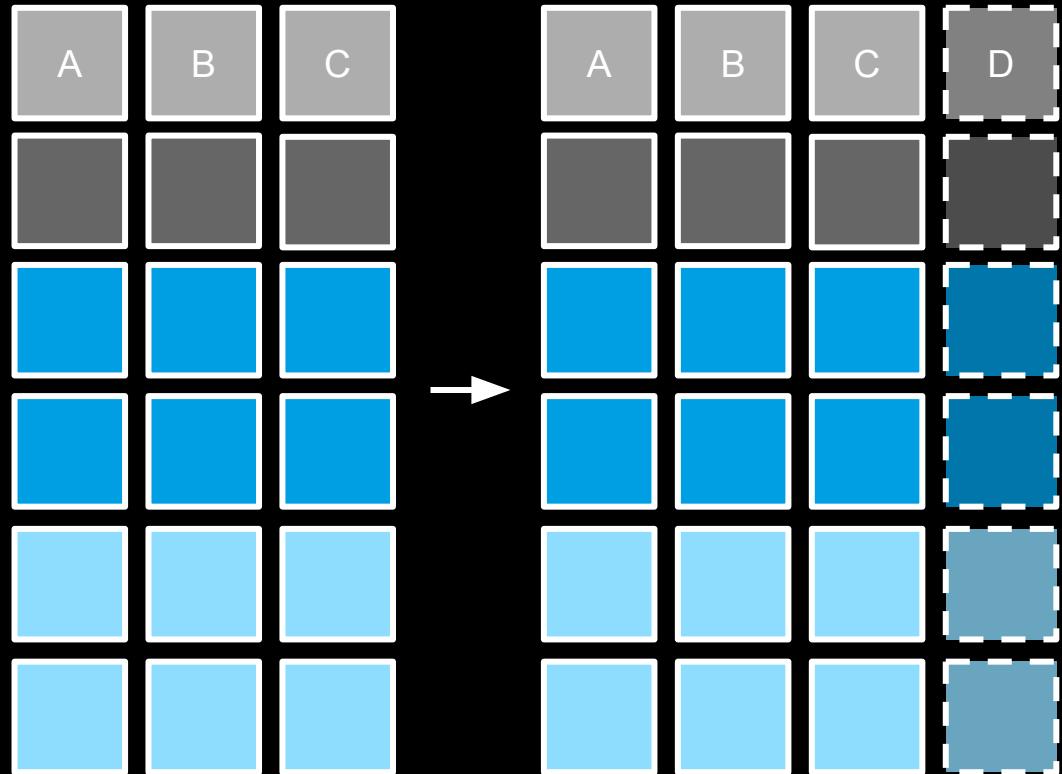
`mutate()`



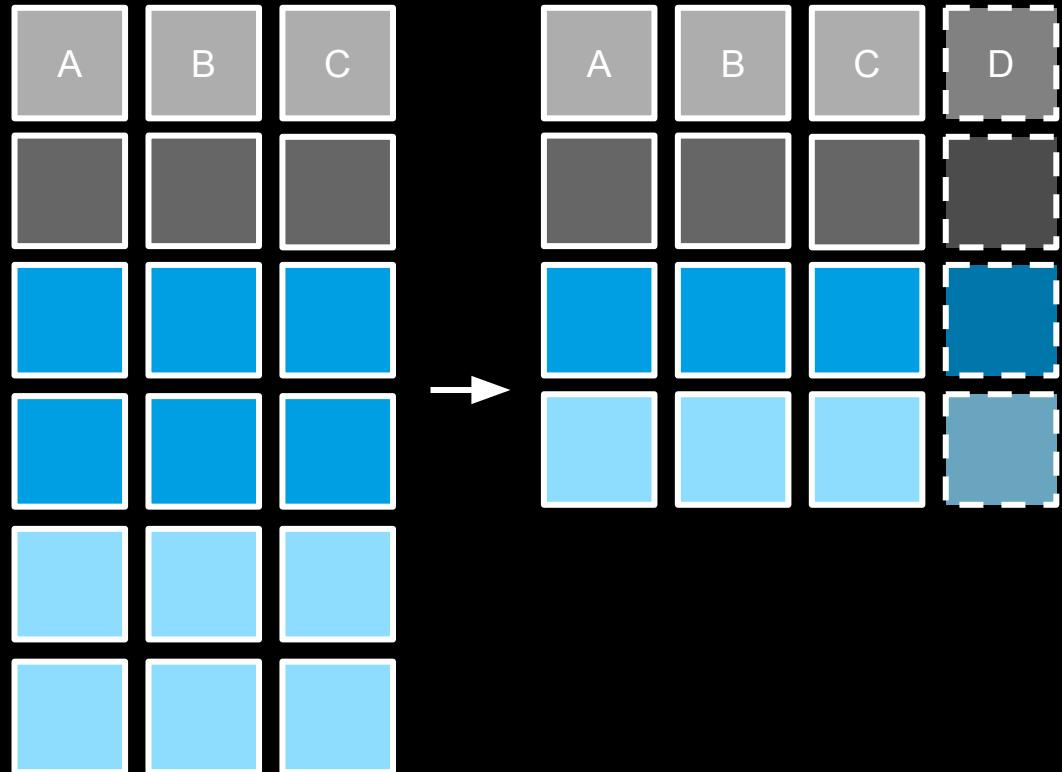
summarize()



group_by()
+
mutate()



group_by()
+
summarize()



joining data

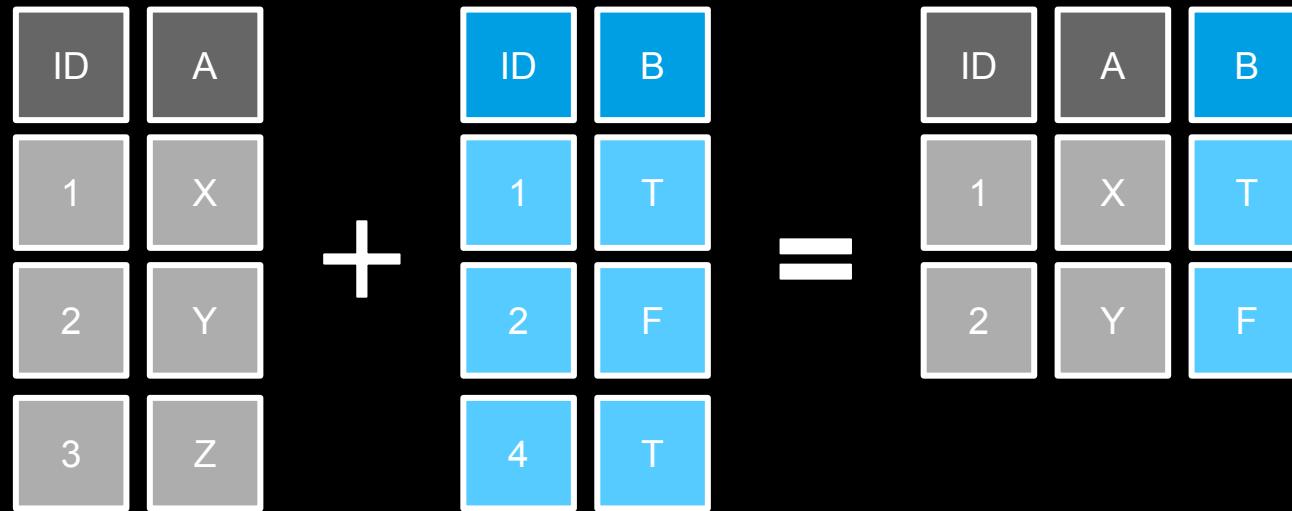
`left_join()`

The diagram illustrates the addition of two 4x2 matrices. The first matrix (left) has rows labeled 1 through 4 and columns labeled ID and A. The second matrix (middle) has rows labeled 1 through 4 and columns labeled ID and B. The result (right) is a third matrix with rows labeled 1 through 4 and columns labeled ID, A, and B. The result matrix contains the following values:

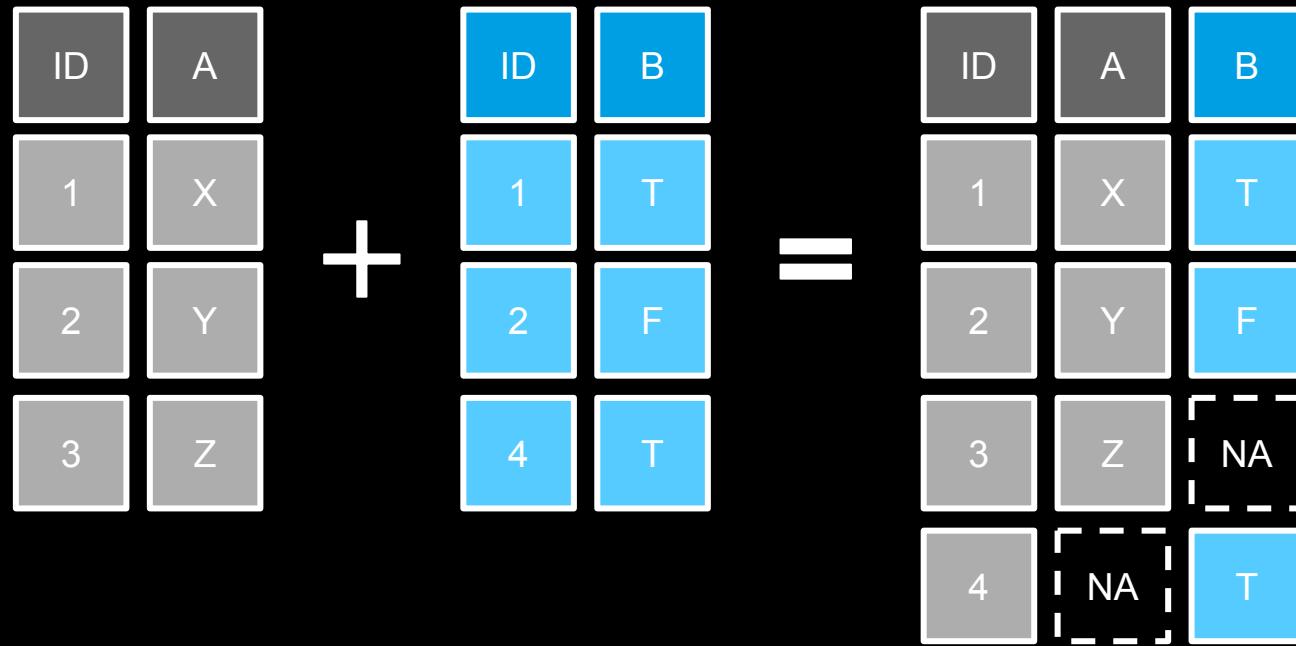
	ID	A	B
1	1	X	T
2	2	Y	F
3	3	Z	NA

The addition is indicated by a plus sign between the first two matrices and an equals sign between the second matrix and the result matrix.

inner_join()



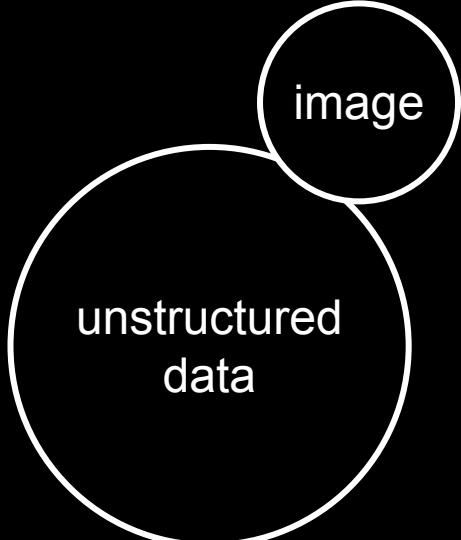
`full_join()`



UNSTRUCTURED DATA

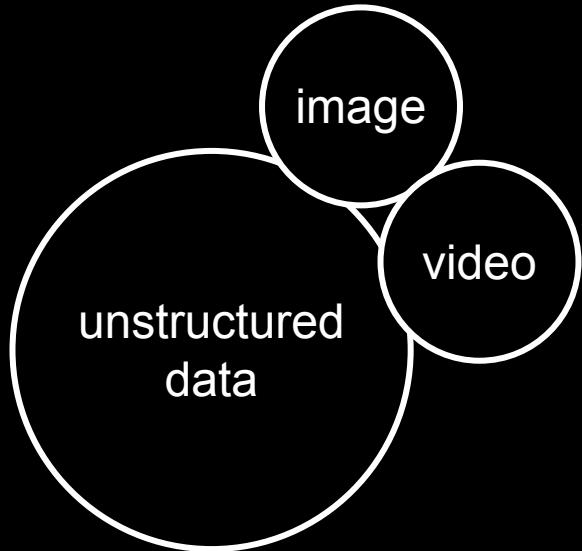


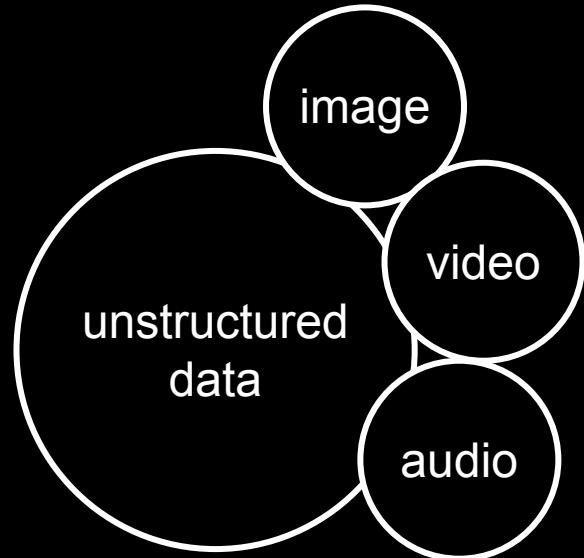
unstructured
data

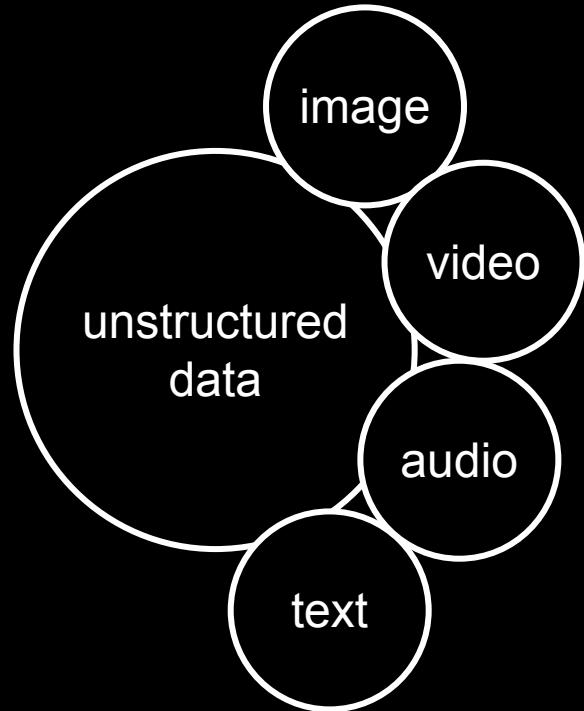


image

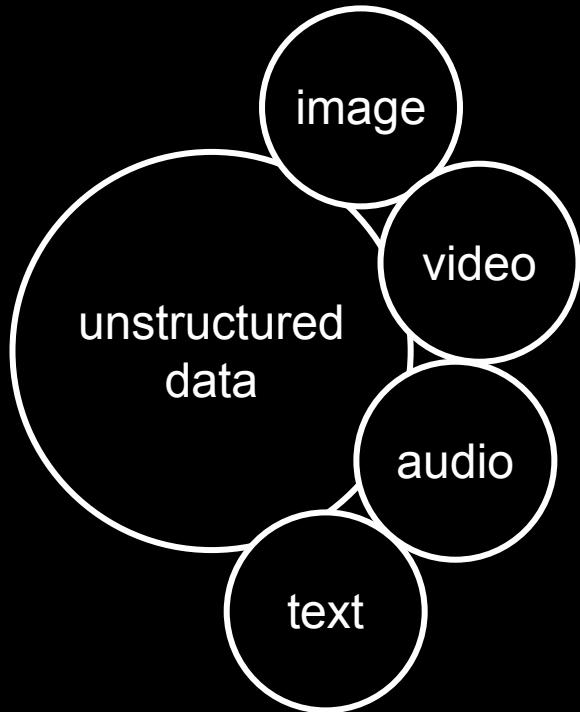
unstructured
data



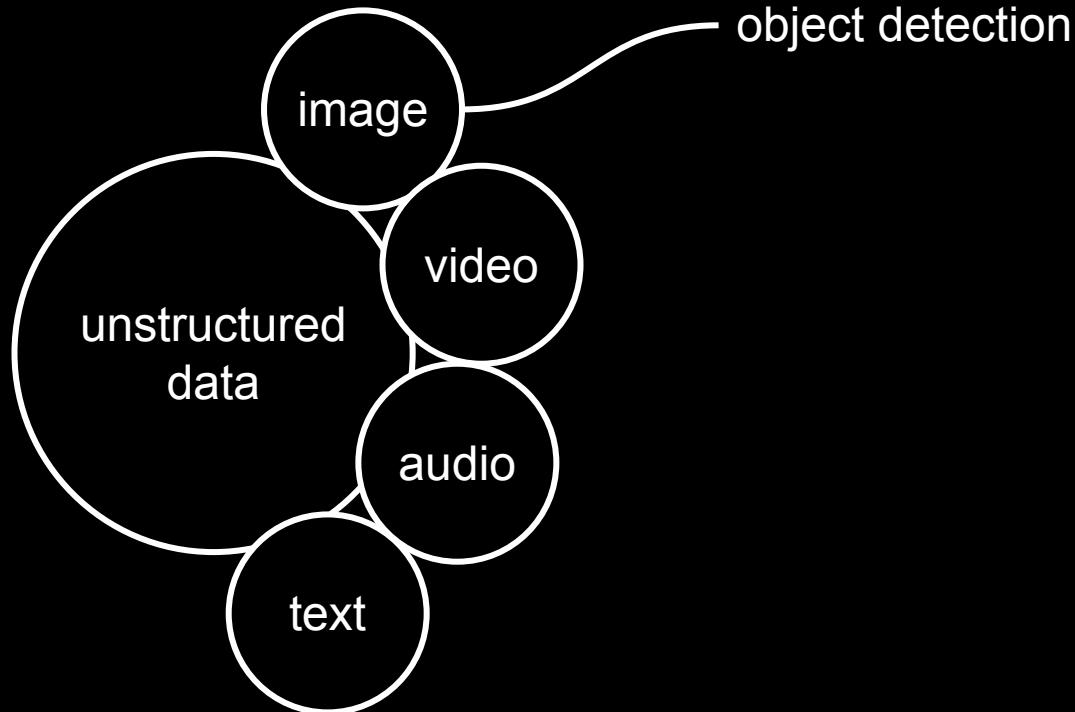




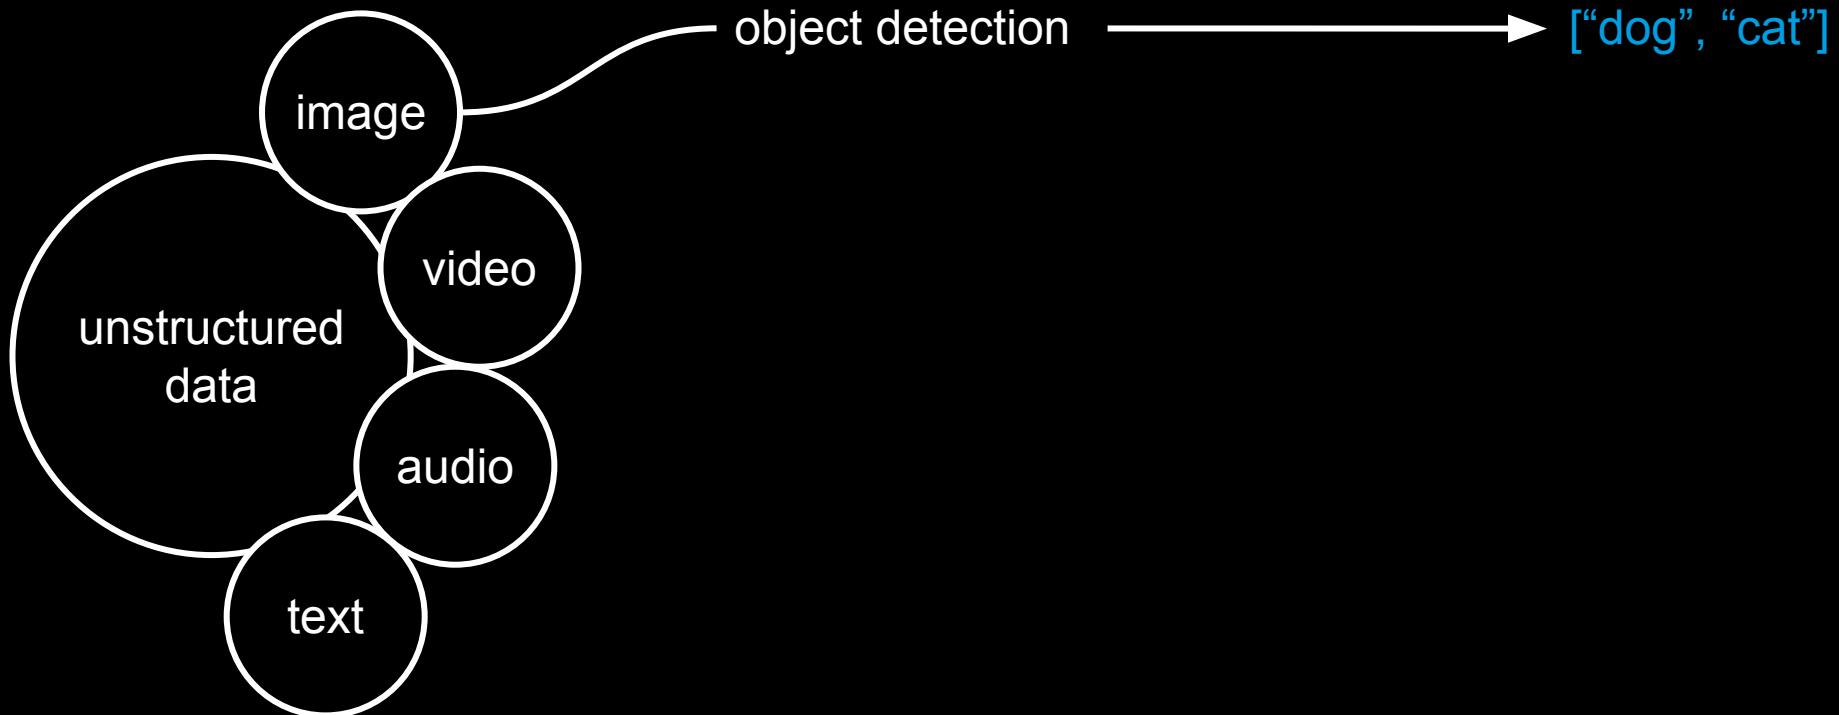
no handles to grab



no handles to grab



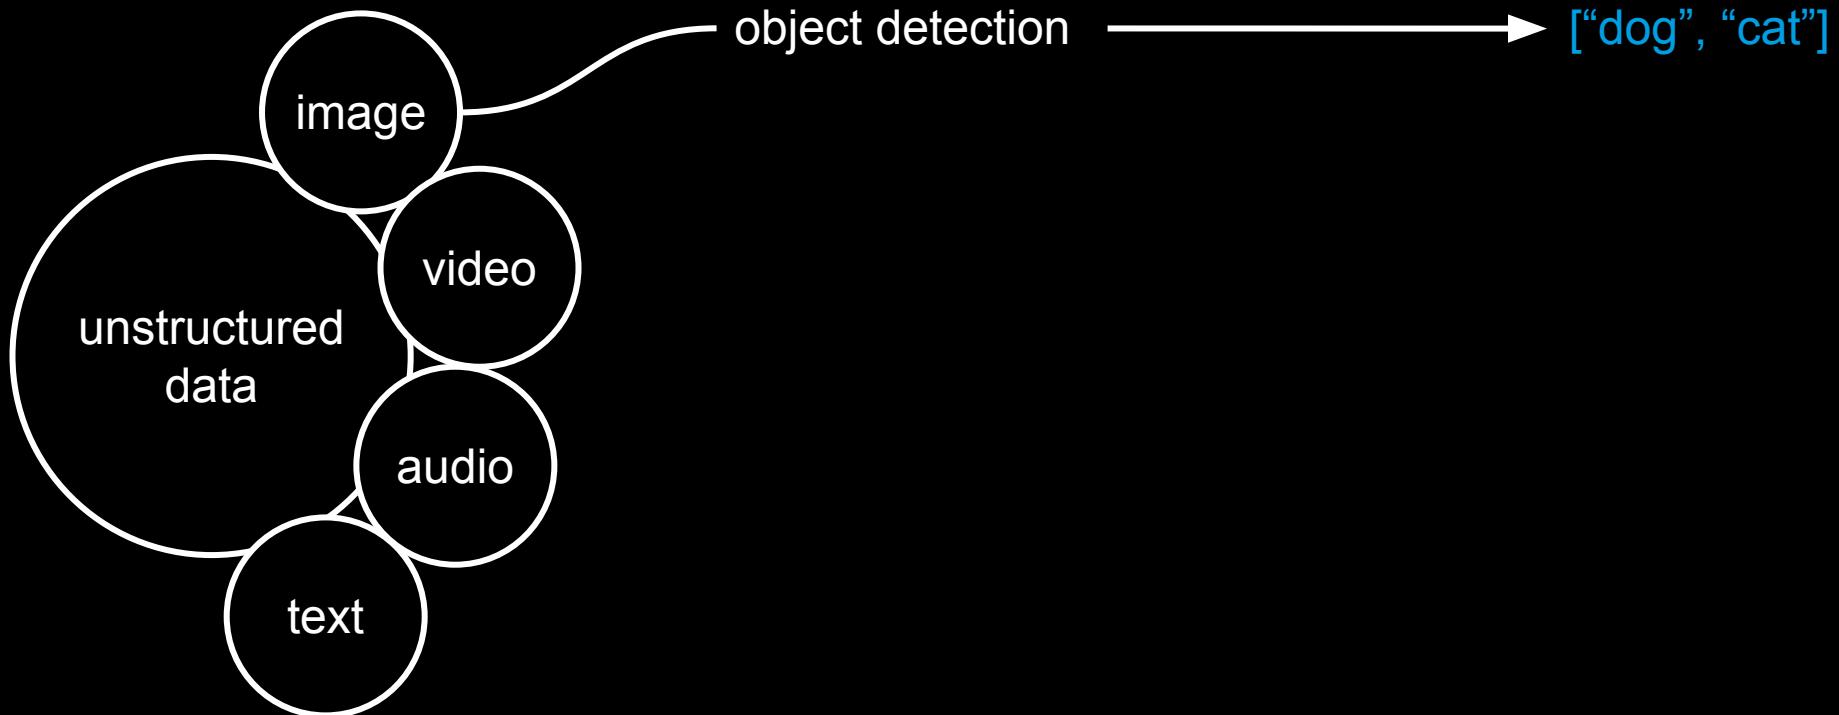
no handles to grab



no handles to grab

algorithm

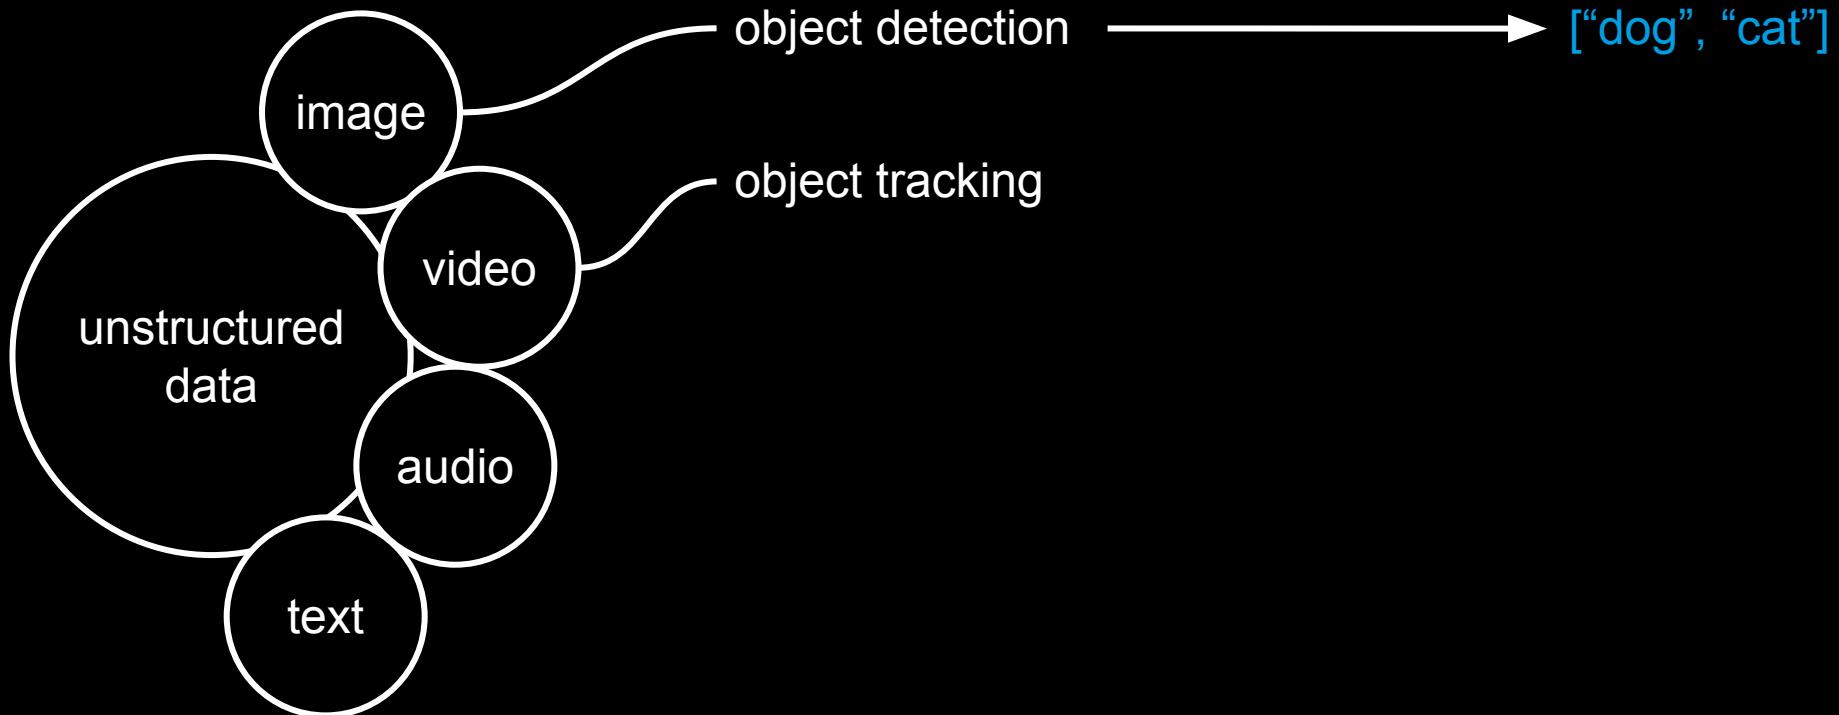
extracted, structured information



no handles to grab

algorithm

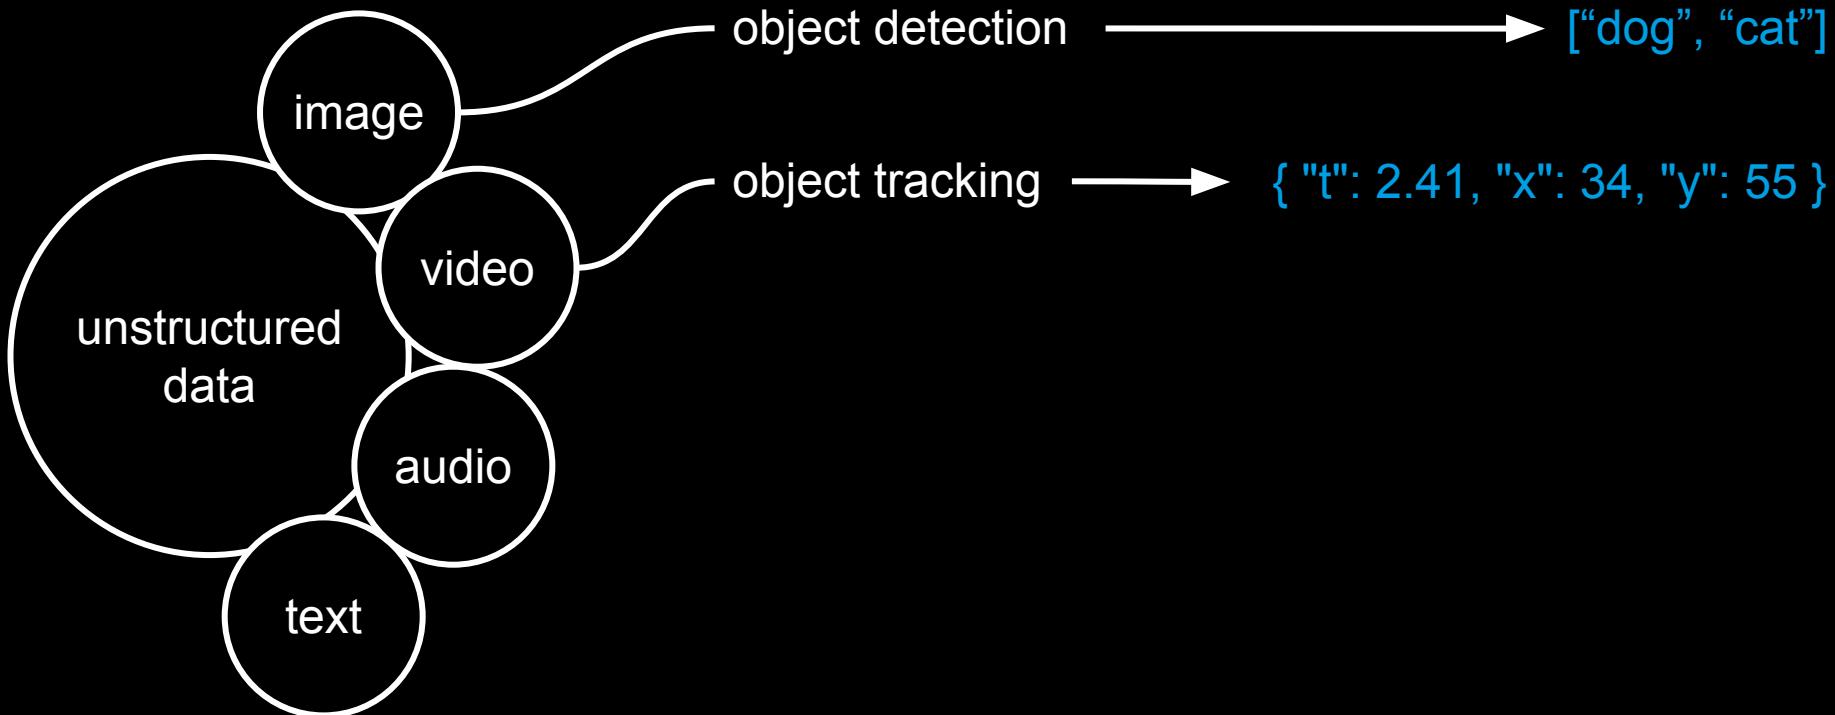
extracted, structured information



no handles to grab

algorithm

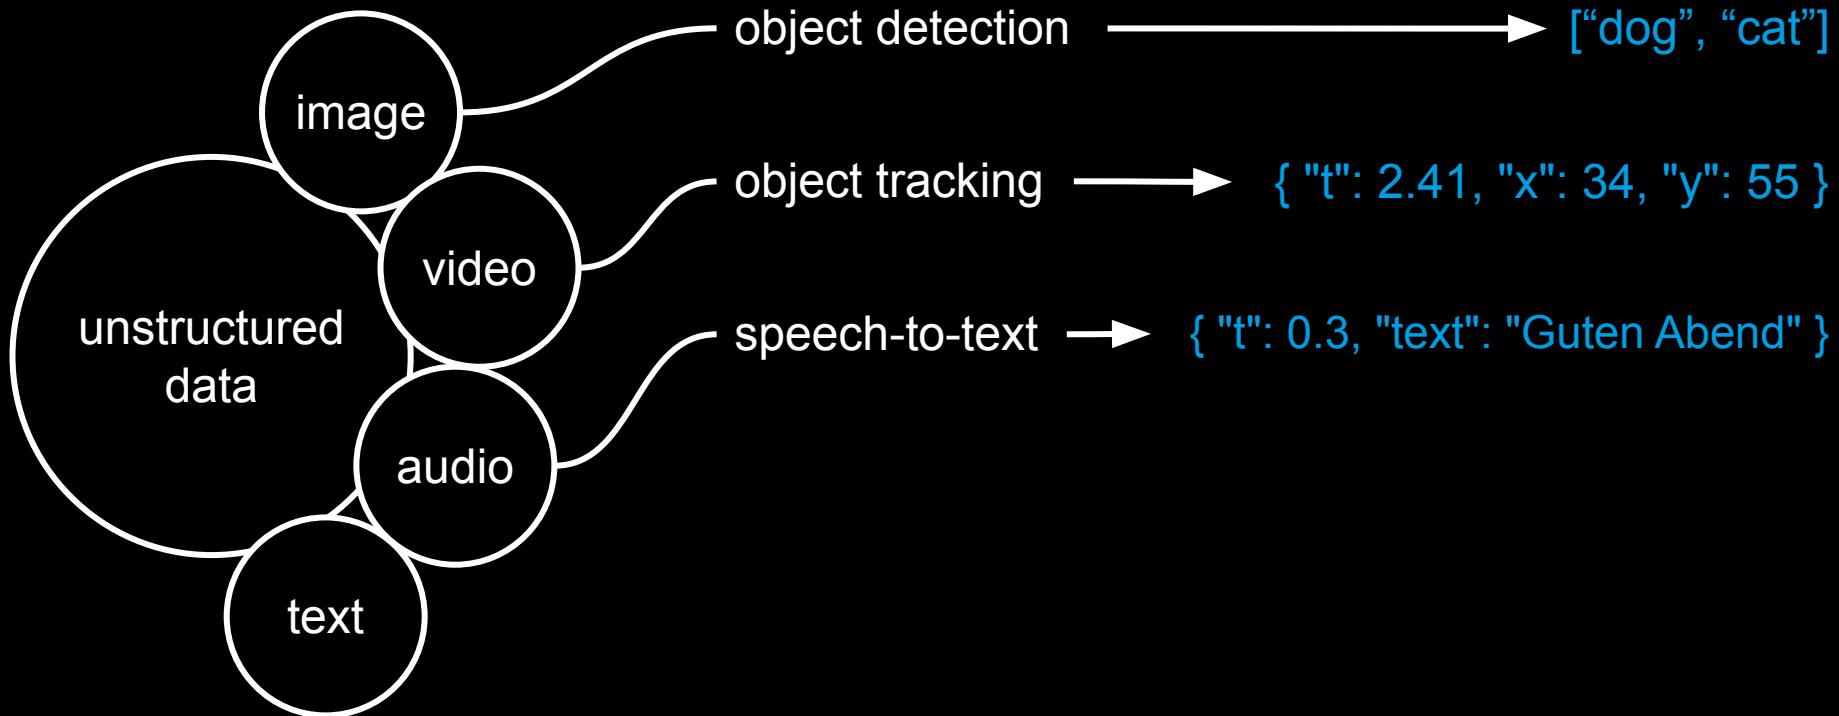
extracted, structured information



no handles to grab

algorithm

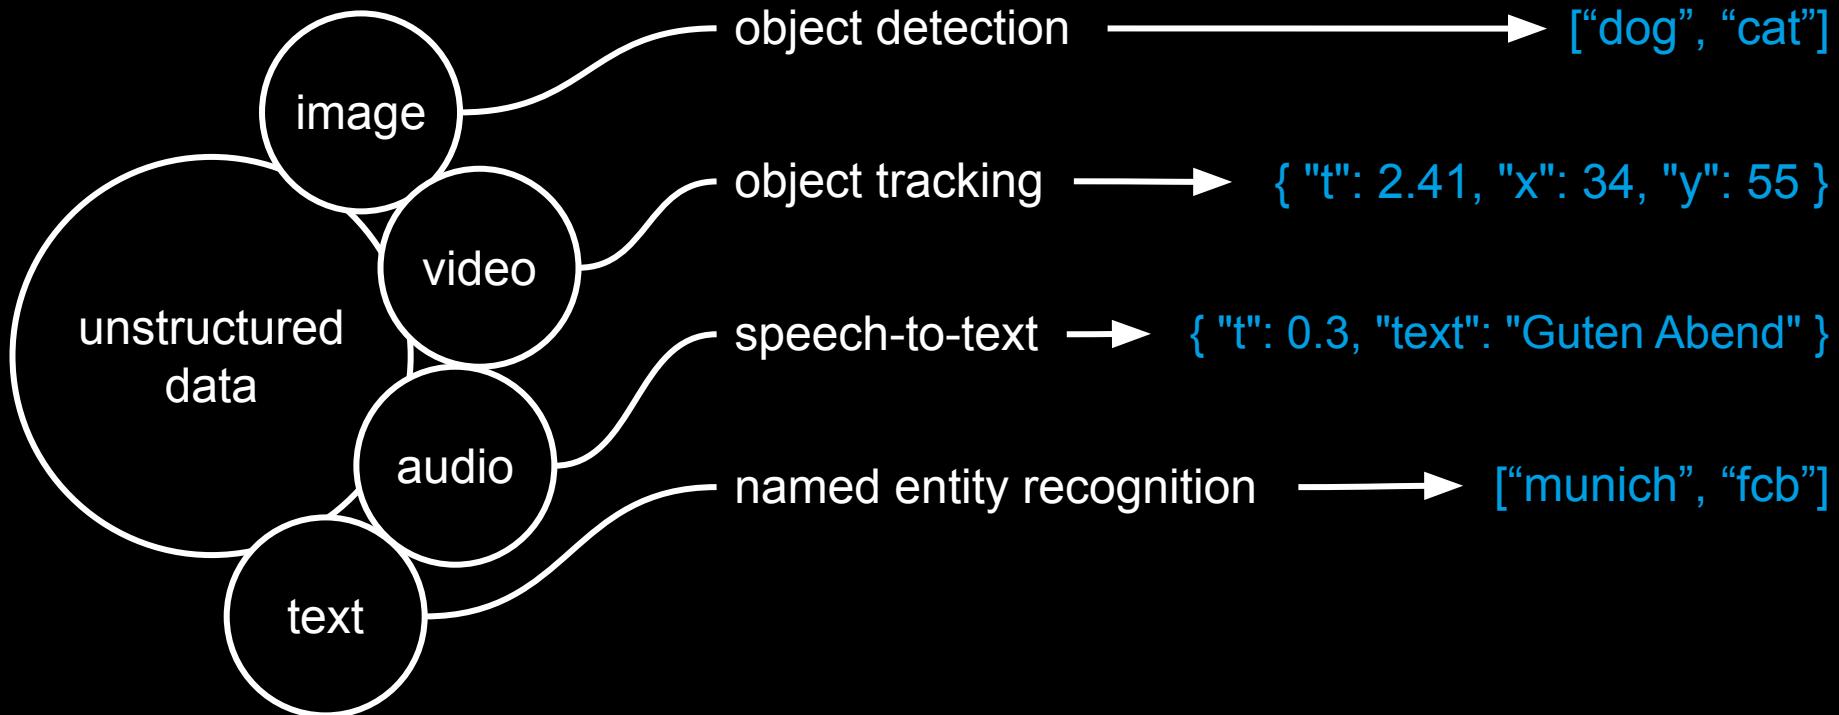
extracted, structured information



no handles to grab

algorithm

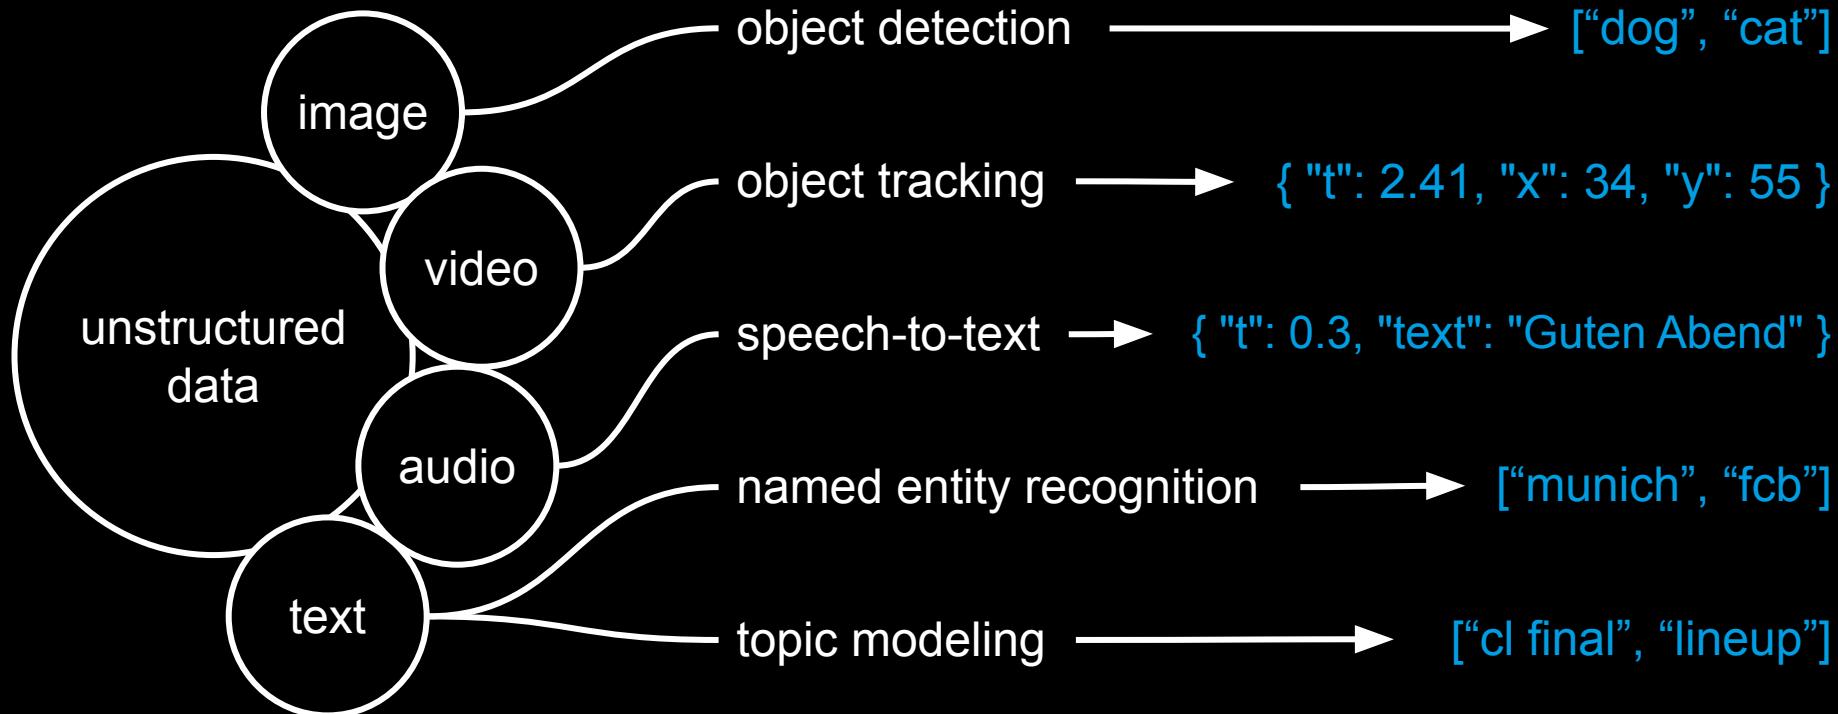
extracted, structured information



no handles to grab

algorithm

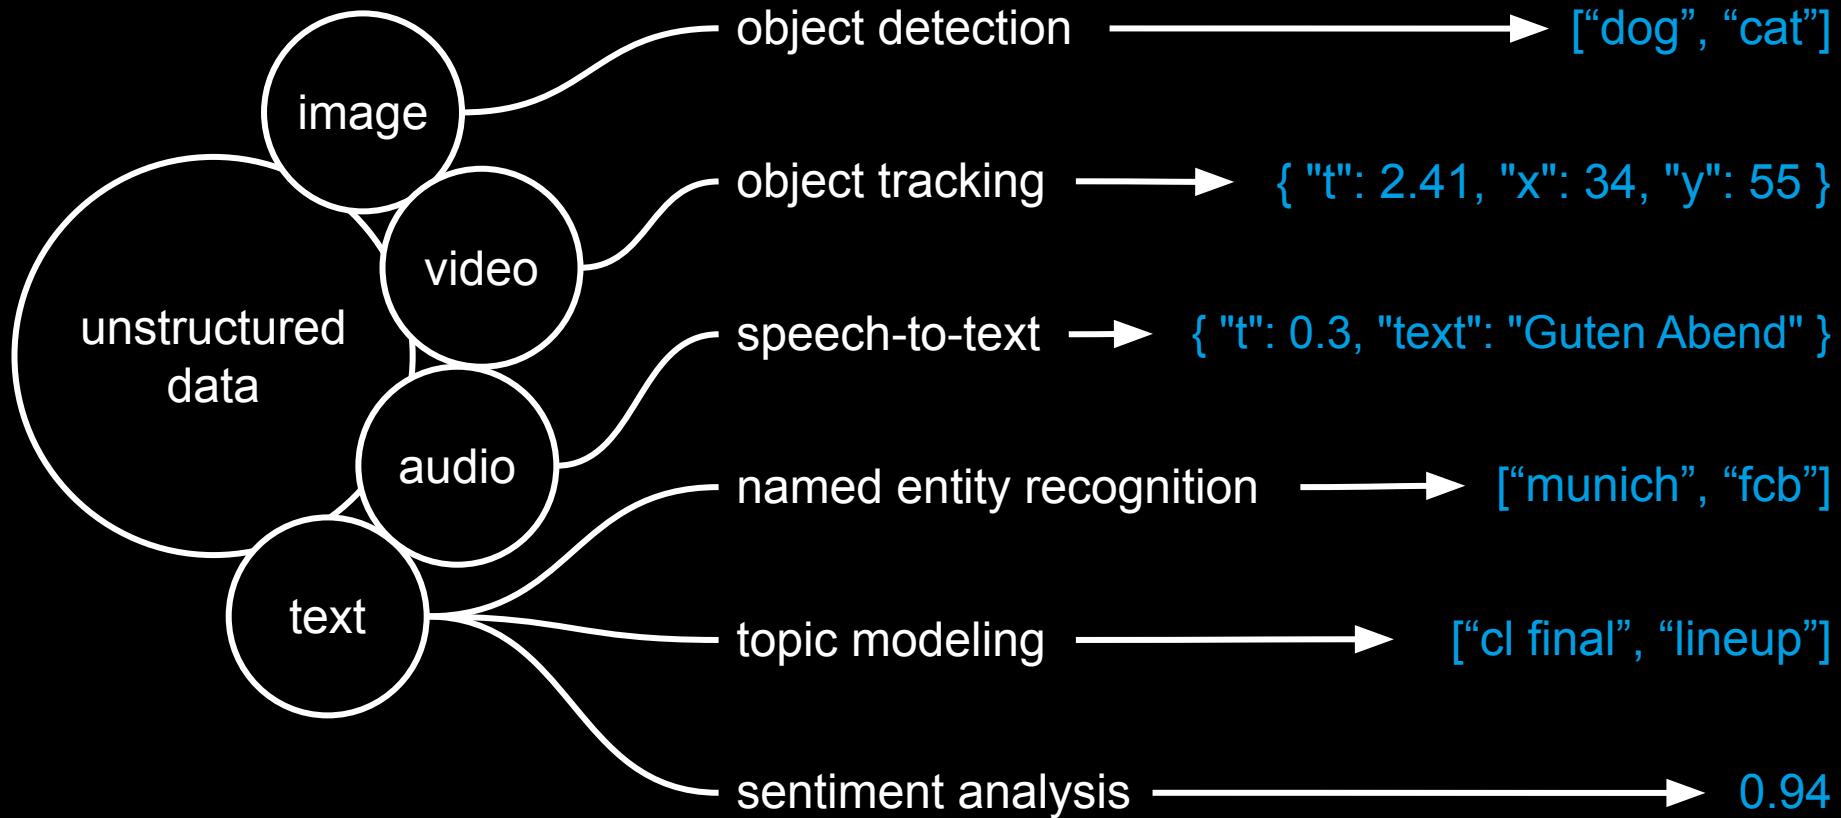
extracted, structured information



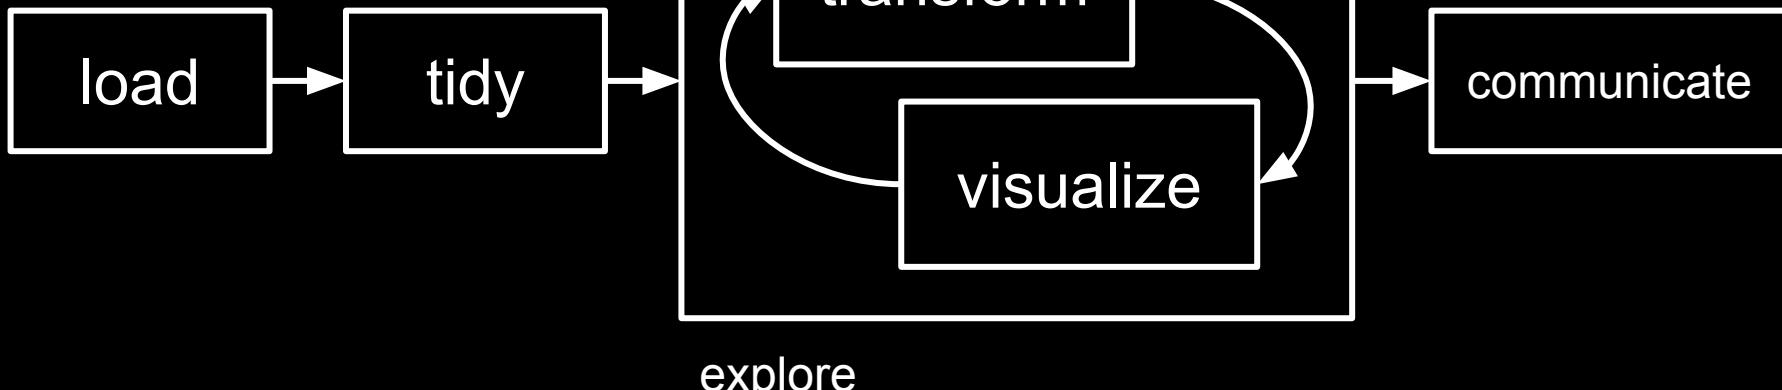
no handles to grab

algorithm

extracted, structured information



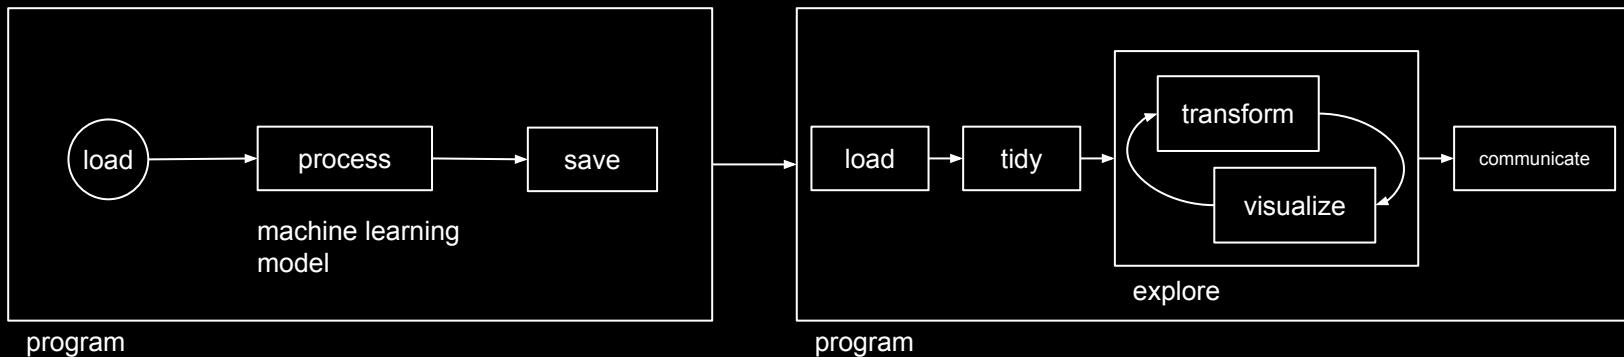
this assumes
structured data



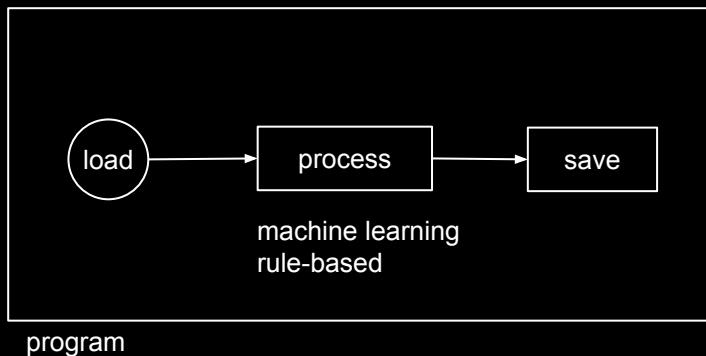
program

pre-process
unstructured data

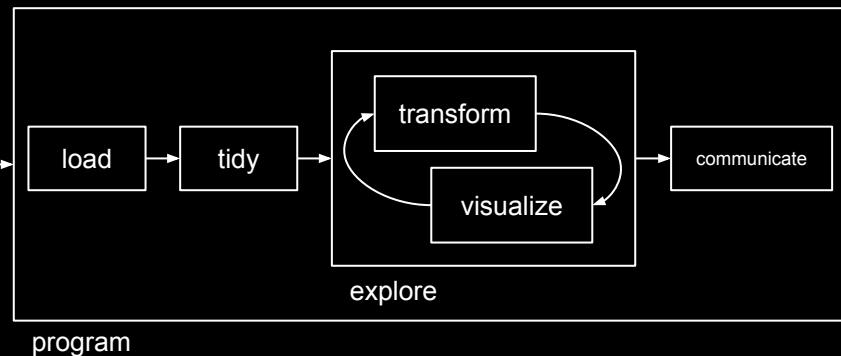
exploratory data
analysis



pre-process unstructured data

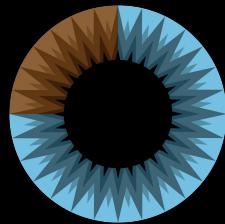


exploratory data analysis



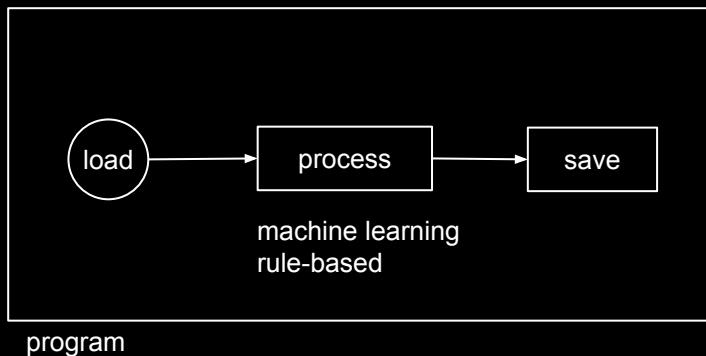
MACHINE LEARNING

Highly recommended for
background information

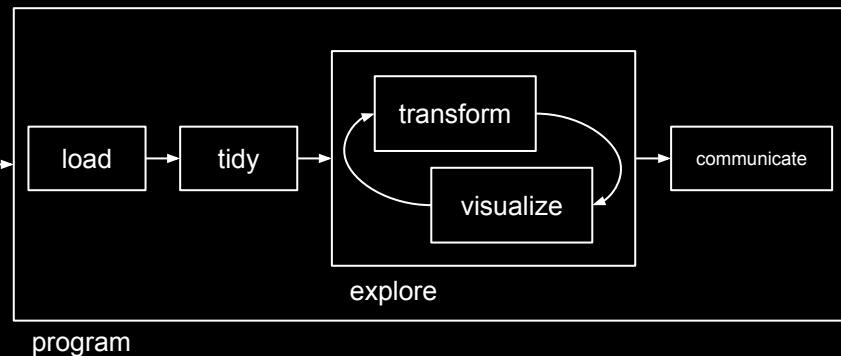


3Blue1Brown's YouTube Course on Neural
Networks and Deep Learning

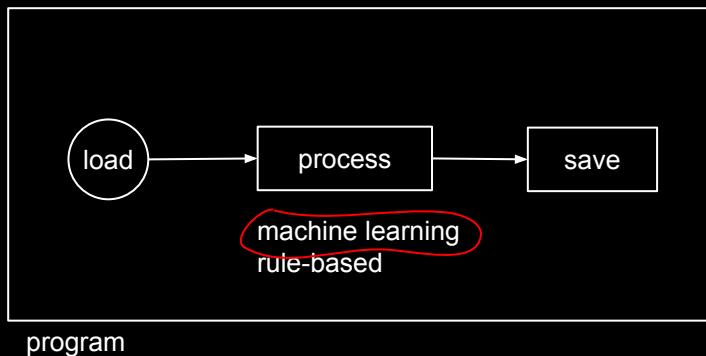
pre-process unstructured data



exploratory data analysis



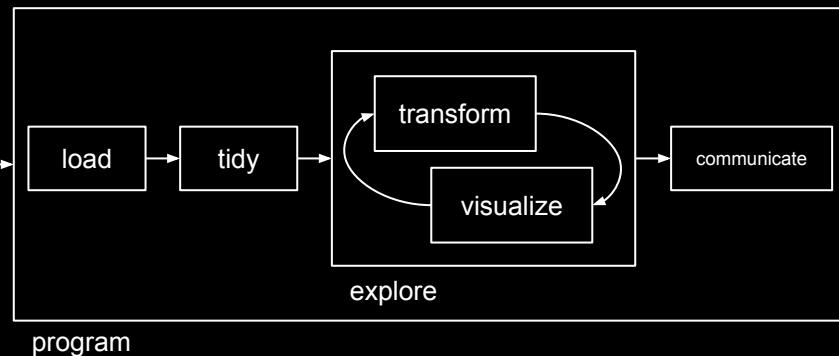
pre-process unstructured data



program



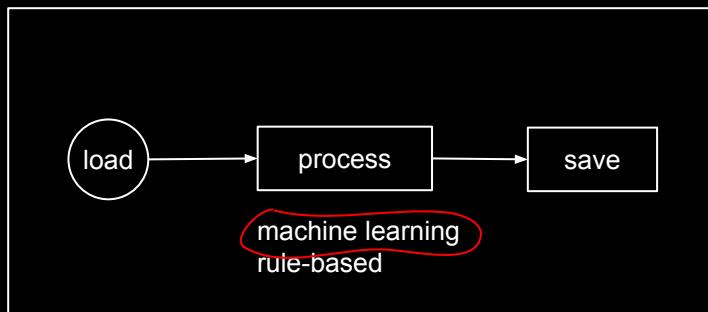
exploratory data analysis



program



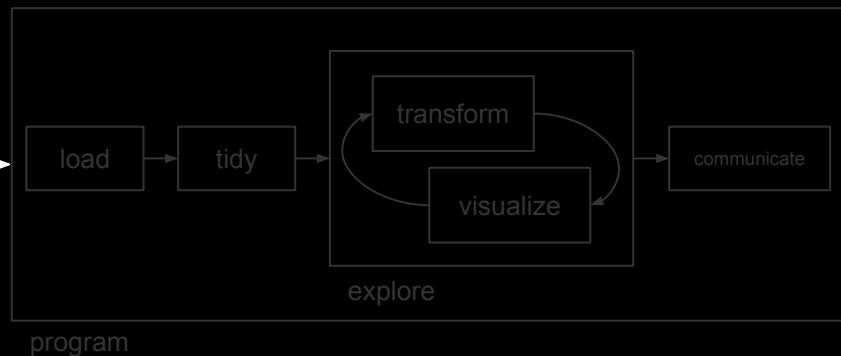
pre-process unstructured data



program



exploratory data analysis

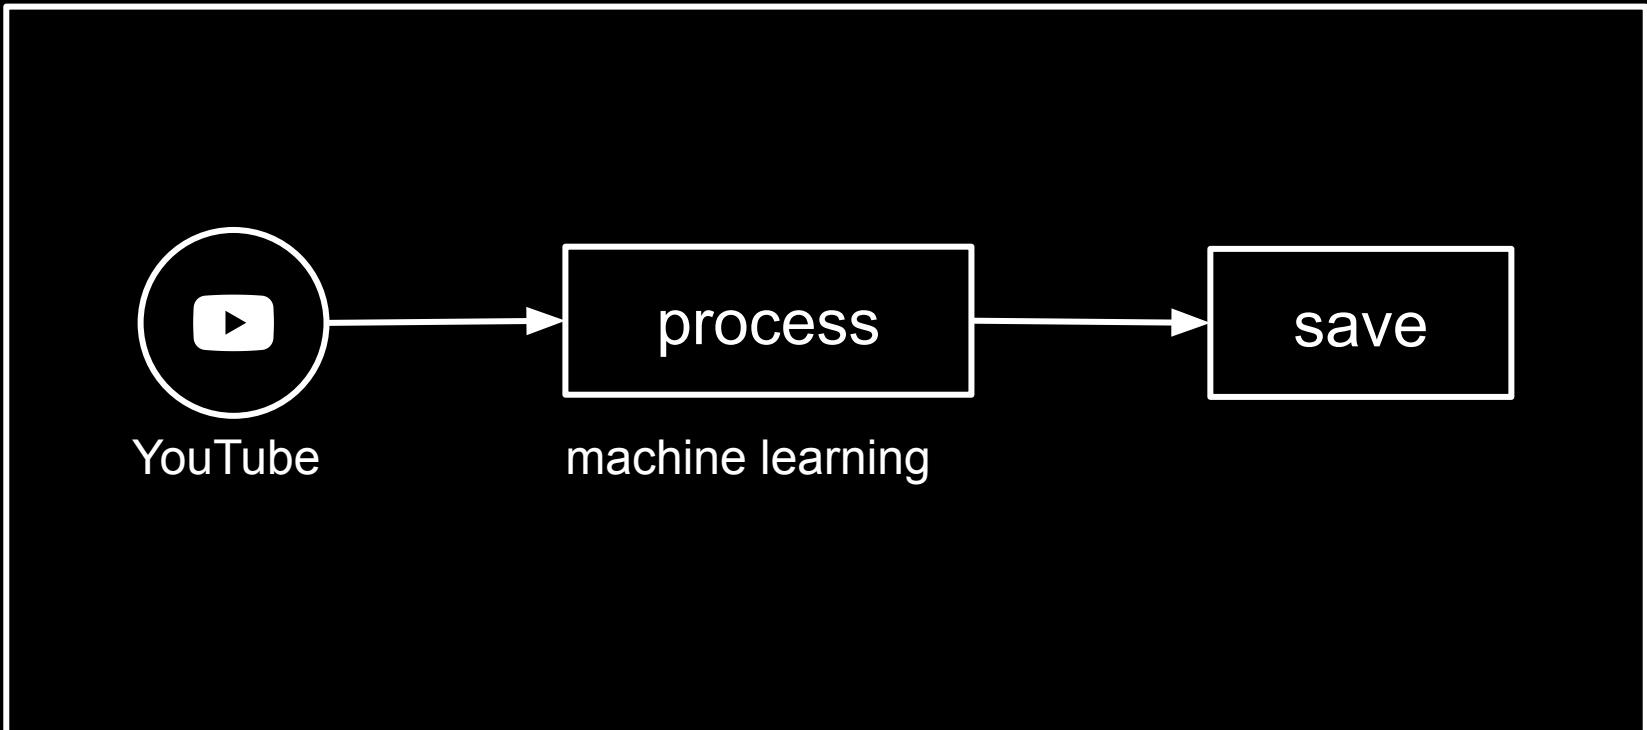


program

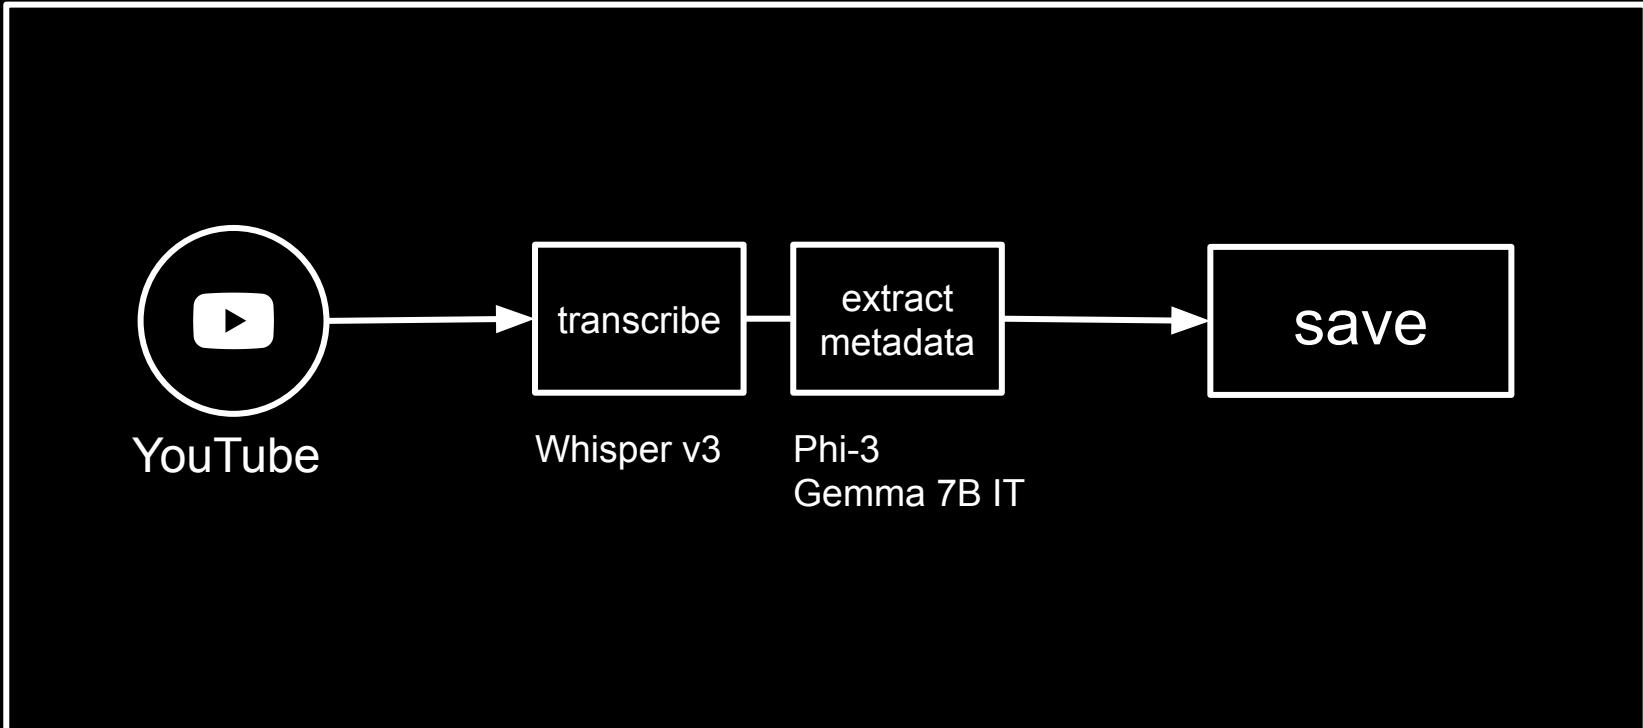




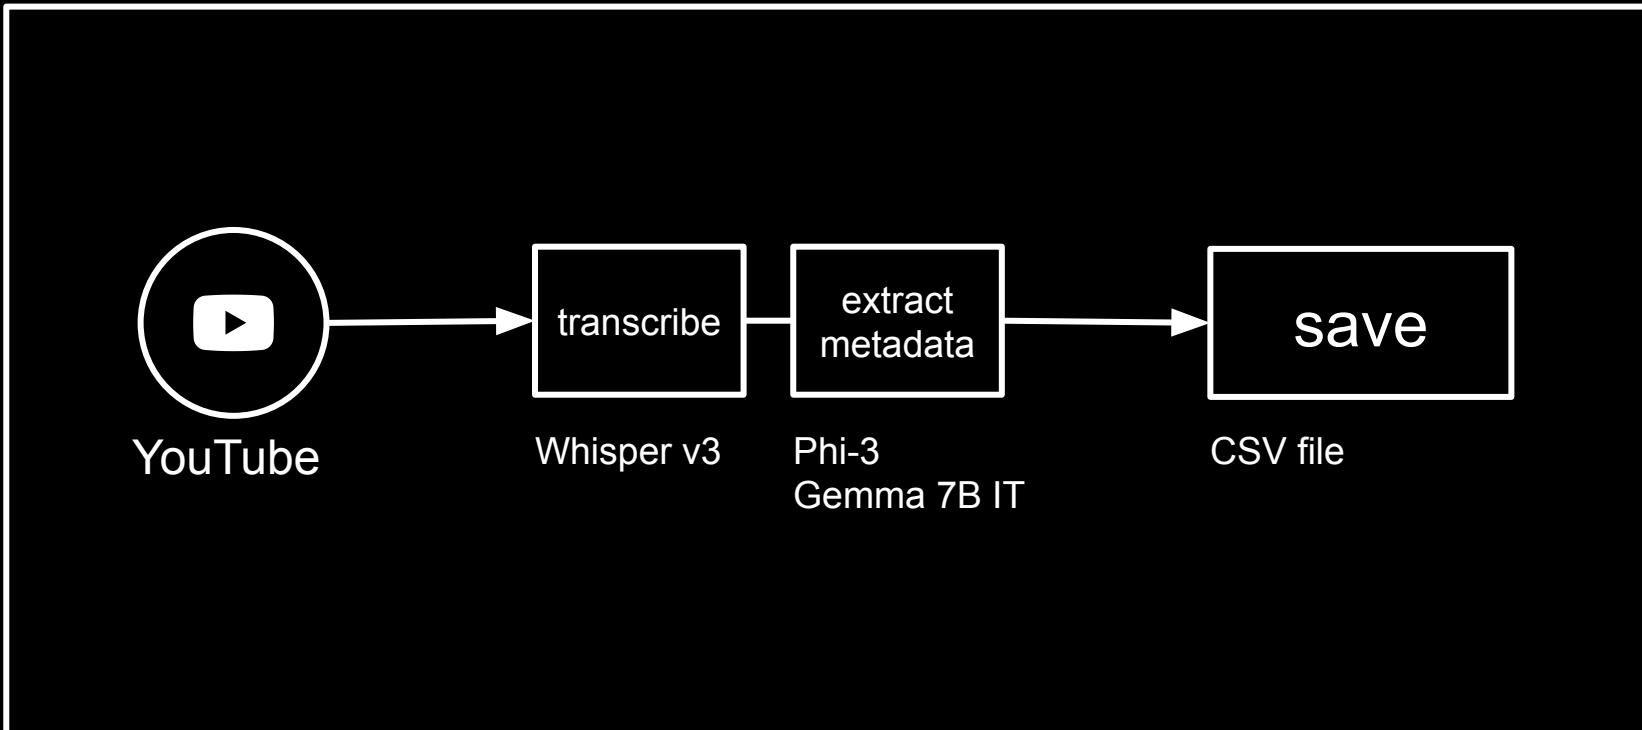
program



program



program



program

YouTube API

Whisper v3

<https://arxiv.org/abs/2212.04356>



<https://huggingface.co/openai/whisper-large-v3>



Large Language Models (LLM)

what has been said so far?
(prompt / context)

what has been said so far?
(prompt / context)



prediction of next token based on
learnt probability distribution

what has been said so far?
(prompt / context)



prediction of next token based on
learnt probability distribution

+

(randomness)

what has been said so far?
(prompt / context)



prediction of next token based on
learnt probability distribution

+

(randomness)

+

(filter)

(discriminating, insulting content)

what has been said so far?
(prompt / context)



prediction of next token based on
learnt probability distribution

+

(randomness)

+

(filter)

(discriminating, insulting content)



next word (*token*)

what has been said so far?
(prompt / context)



prediction of next token based on
learnt probability distribution



next word (*token*)

+
(randomness)
+
(filter)
(discriminating, insulting content)



PROMPTING

<https://www.promptingguide.ai/>



elements of a prompt

```
<instruction>  
<context>  
<input data>  
<output indicator>
```

elements of a prompt

<instruction>

<context>

<input data>

<output indicator>

example prompt

Explain the binary number system.

elements of a prompt

<instruction>

<context>

<input data>

<output indicator>

example prompt

Explain the binary number system.

start simple

elements of a prompt

<instruction>
<context>
<input data>
<output indicator>

example prompt

You are a friendly tutor and your task is to explain complex concepts as simple as possible.

Explain the binary number system.

elements of a prompt

<instruction>
<context>
<input data>
<output indicator>

example prompt

You are a friendly tutor and your task is to explain complex concepts as simple as possible.

Your answers are never longer than 10 sentences.

Explain the binary number system.

ZERO-SHOT PROMPTING

elements of a prompt

<instruction>

<context>

<input data>

<output indicator>

example prompt

Classify the text into neutral,
negative or positive.

Text: "What a great dinner!"

Sentiment:

elements of a prompt

<instruction>
<context>
<input data>
<output indicator>

example prompt

Classify the text into neutral,
negative or positive.

Text: "What a great dinner!"

Sentiment:

this will be replaced with
data later...

FEW-SHOT PROMPTING

IN-CONTEXT LEARNING

examples in the context to learn from

Extract all references to countries and their continent in the following text using the format from the examples below.

Example 1: "They played the team called 'Die Mannschaft' in the world cup final"

Correct answer: Germany, Europe

Example 2: "The Three Lions once again lost to Germany in a semi final"

Correct answer: England, Europe, Germany, Europe

Text: "The Selecao was destroyed 1:7 by the DFB selection in their home stadium."

Answer:

examples in the context to learn from

Extract all references to countries and their continent in the following text using the format from the examples below.

Example 1: "They played the team called 'Die Mannschaft' in the world cup final"

Correct answer: Germany, Europe

Example 2: "The Three Lions once again lost to Germany in a semi final"

Correct answer: England, Europe, Germany, Europe

Text: "The Selecao was destroyed 1:7 by the DFB selection in their home stadium."

Answer:

more prompting strategies

chain-of-thought (CoT)

self-consistency

generate knowledge prompting

prompt chaining (subtasks)

tree-of-thoughts (ToT)

retrieval-augmented-generation (RAG)

...

Phi-3

<https://arxiv.org/abs/2404.14219>



https://huggingface.co/microsoft/Phi_3-mini-128k-instruct

<https://huggingface.co/microsoft/Phi-3-medium-128k-instruct>

Gemma 2B / 7B Instruct

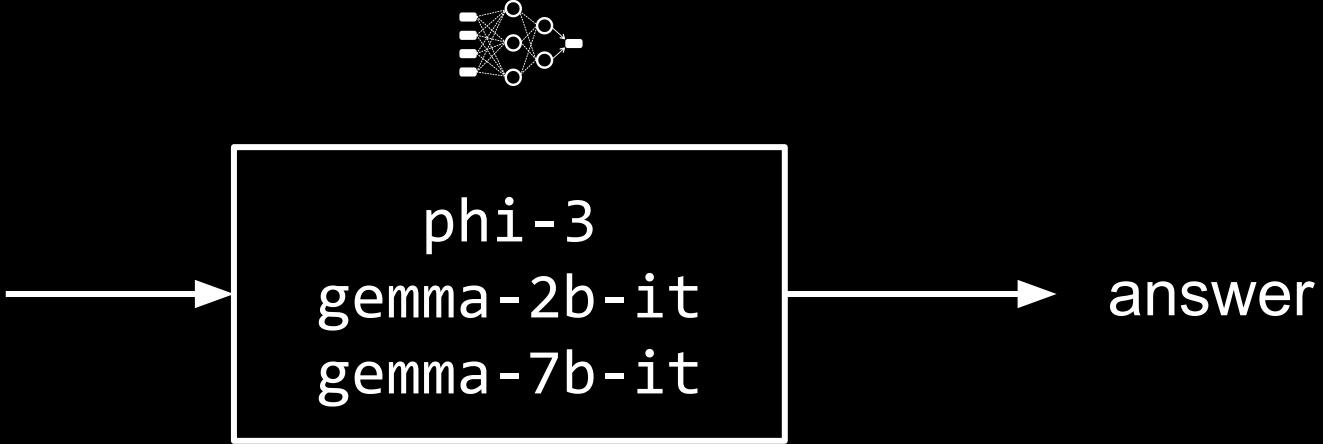
<https://arxiv.org/abs/2403.08295>



<https://huggingface.co/google/gemma-2b-it>

<https://huggingface.co/google/gemma-7b-it>

prompt +
context



OpenAI's GPT-4.1 family

<https://platform.openai.com/docs/models>

 GPT-4.1 nano (Default) ⚙️

Fastest, most cost-effective GPT-4.1 model

Compare Try in Playground

INTELLIGENCE	SPEED	PRICE	INPUT	OUTPUT
● ● Average	⚡ ⚡ ⚡ ⚡ ⚡ Very fast	\$0.1 - \$0.4 Input · Output	Text, image	Text

GPT-4.1 nano is the fastest, most cost-effective GPT-4.1 model.

1,047,576 context window
32,768 max output tokens
Jun 01, 2024 knowledge cutoff

Pricing

Pricing is based on the number of tokens used. For tool-specific models, like search and computer use, there's a fee per tool call. See details in the [pricing page](#).

Text tokens

Per 1M tokens • Batch API price

Input \$0.10	Cached input \$0.025	Output \$0.40
-----------------	-------------------------	------------------

<https://platform.openai.com/docs/models/gpt-4.1-nano>

1 mio. token →
~2500 pages

The screenshot shows the GPT-4.1 nano model page on the OpenAI platform. At the top, it says "GPT-4.1 nano (Default)" and "Fastest, most cost-effective GPT-4.1 model". Below this are sections for Intelligence (Average), Speed (Very fast), Price (\$0.1 - \$0.4), Input (Text, image), and Output (Text). A large orange arrow points from the text above to the "1,047,576 context window" information, which is highlighted with an orange circle. This section also includes "32,768 max output tokens" and a knowledge cutoff date of "Jun 01, 2024".

INTELLIGENCE: Average

SPEED: Very fast

PRICE: \$0.1 - \$0.4

INPUT: Text, image

OUTPUT: Text

GPT-4.1 nano is the fastest, most cost-effective GPT-4.1 model.

1,047,576 context window
32,768 max output tokens
Jun 01, 2024 knowledge cutoff

Pricing: Pricing is based on the number of tokens used. For tool-specific models, like search and computer use, there's a fee per tool call. See details in the [pricing page](#).

Text tokens: Input \$0.10, Cached input \$0.025, Output \$0.40

Per 1M tokens • Batch API price

<https://platform.openai.com/docs/models/gpt-4.1-nano>

1 mio. token →
~2500 pages

roughly 10 cents
as input

The screenshot shows the GPT-4.1 nano model page on the OpenAI platform. At the top, it says "GPT-4.1 nano (Default)" and "Fastest, most cost-effective GPT-4.1 model". Below this are sections for "INTELLIGENCE" (Average, represented by two dots), "SPEED" (Very fast, represented by five lightning bolts), and "PRICE" (\$0.1 - \$0.4, Input + Output). On the right, there are "INPUT" and "OUTPUT" options for "Text, image" and "Text". A callout arrow points from the text "roughly 10 cents as input" to the "Input" price of \$0.10. Another callout arrow points from the text "1 mio. token → ~2500 pages" to the "Context window" of 1,047,576 tokens.

GPT-4.1 nano is the fastest, most cost-effective GPT-4.1 model.

INTELLIGENCE: Average

SPEED: Very fast

PRICE: \$0.1 - \$0.4

INPUT: Text, image

OUTPUT: Text

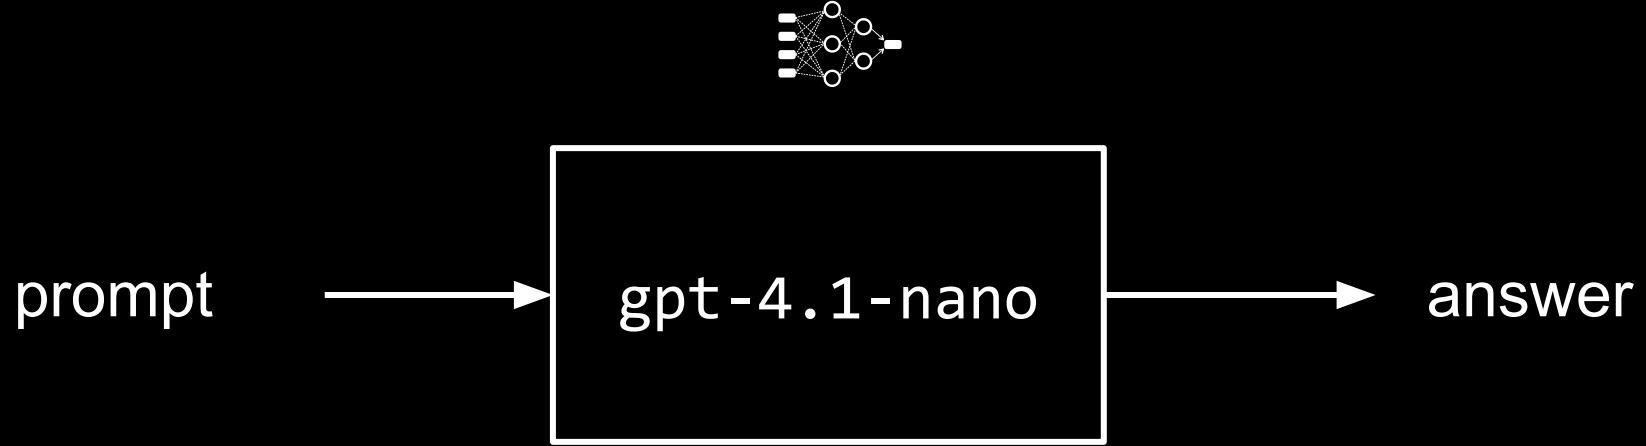
Pricing: Pricing is based on the number of tokens used. For tool-specific models, like search and computer use, there's a fee per tool call. See details in the [pricing page](#).

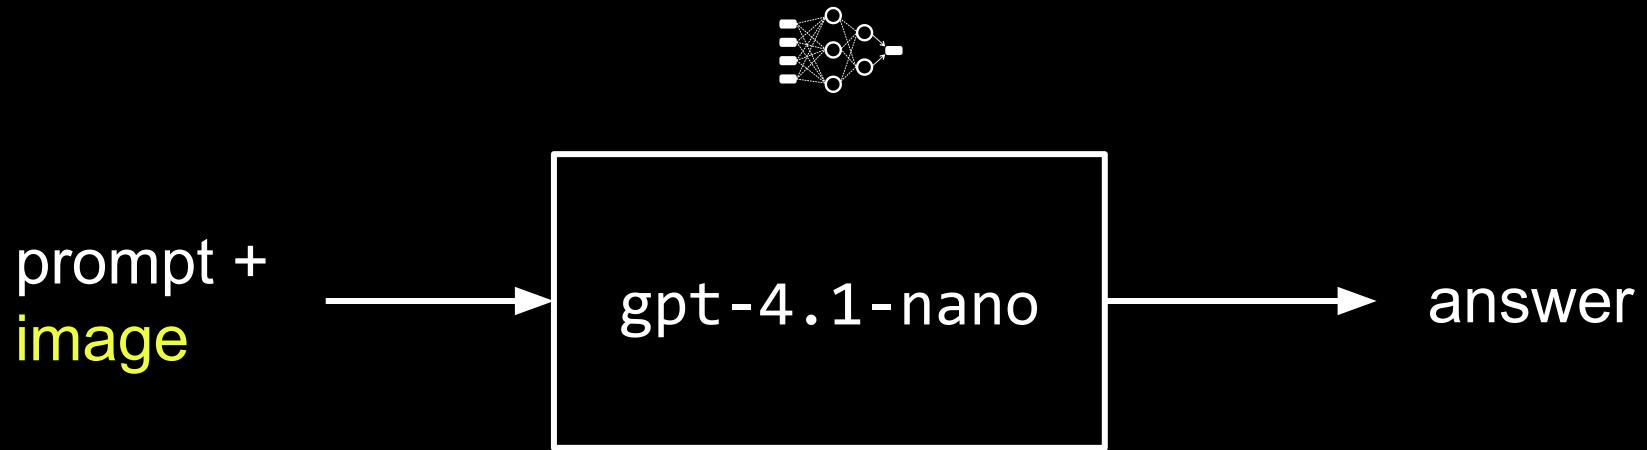
Text tokens: Input \$0.10, Cached input \$0.025, Output \$0.40

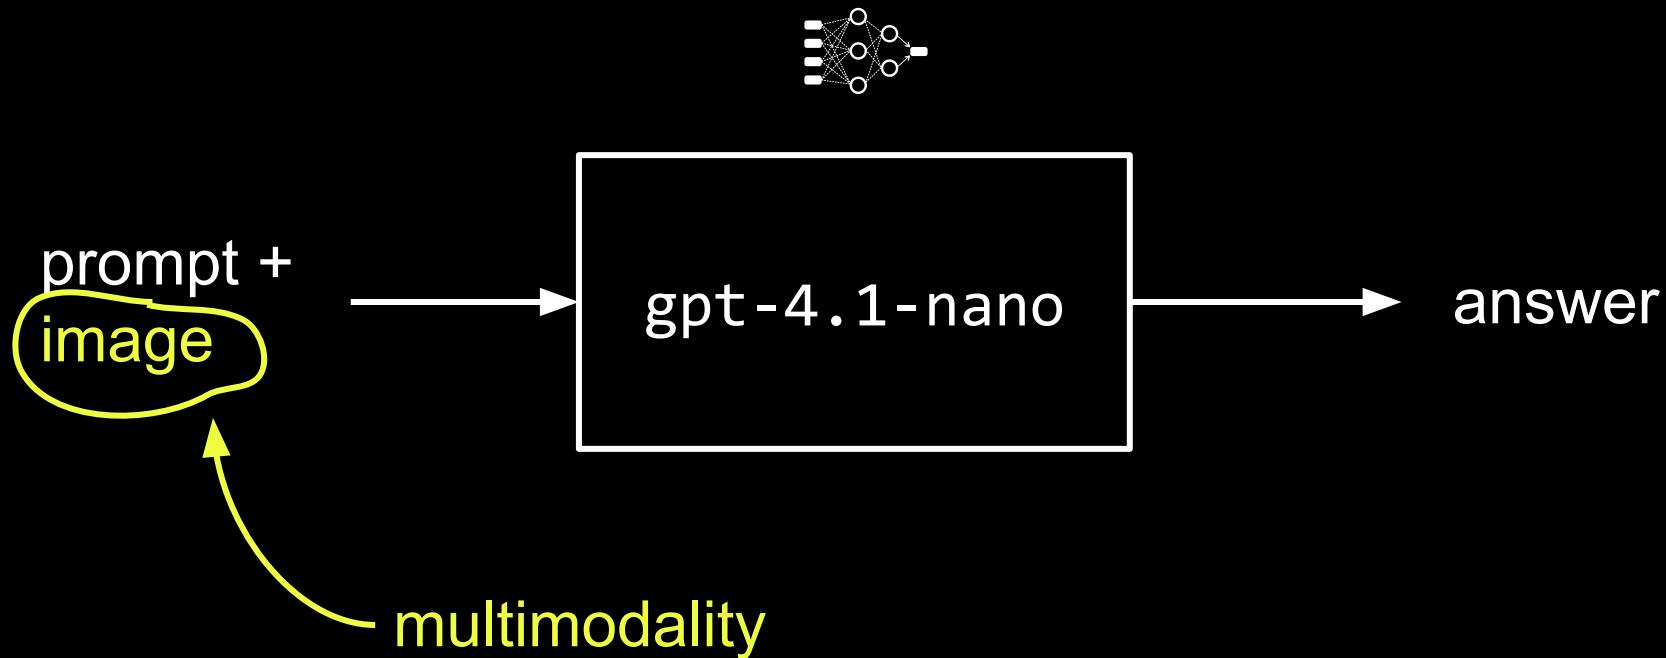
Per 1M tokens • Batch API price

1,047,576 context window
32,768 max output tokens
Jun 01, 2024 knowledge cutoff

<https://platform.openai.com/docs/models/gpt-4.1-nano>



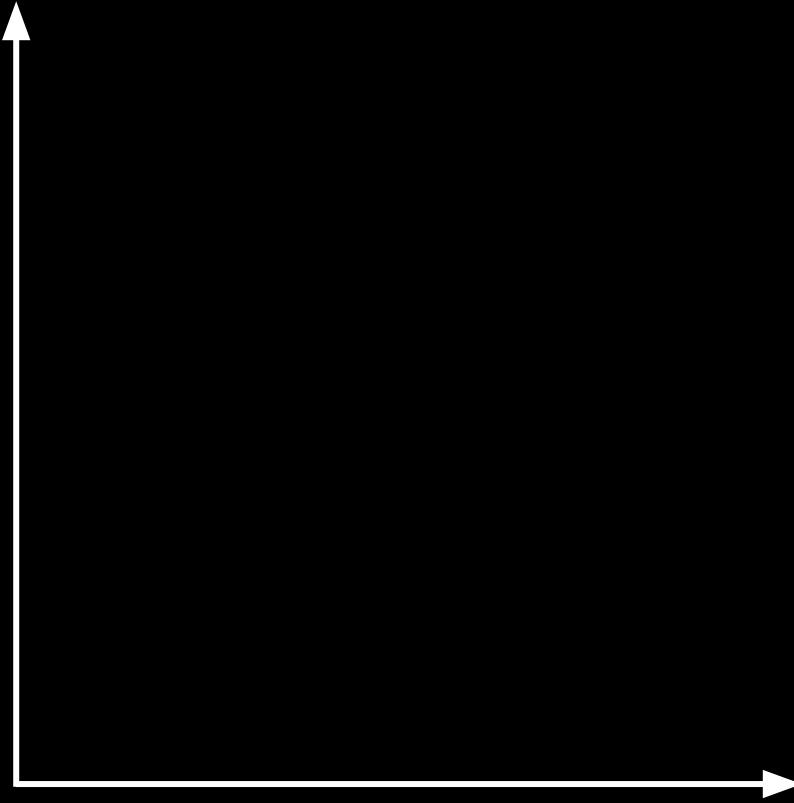


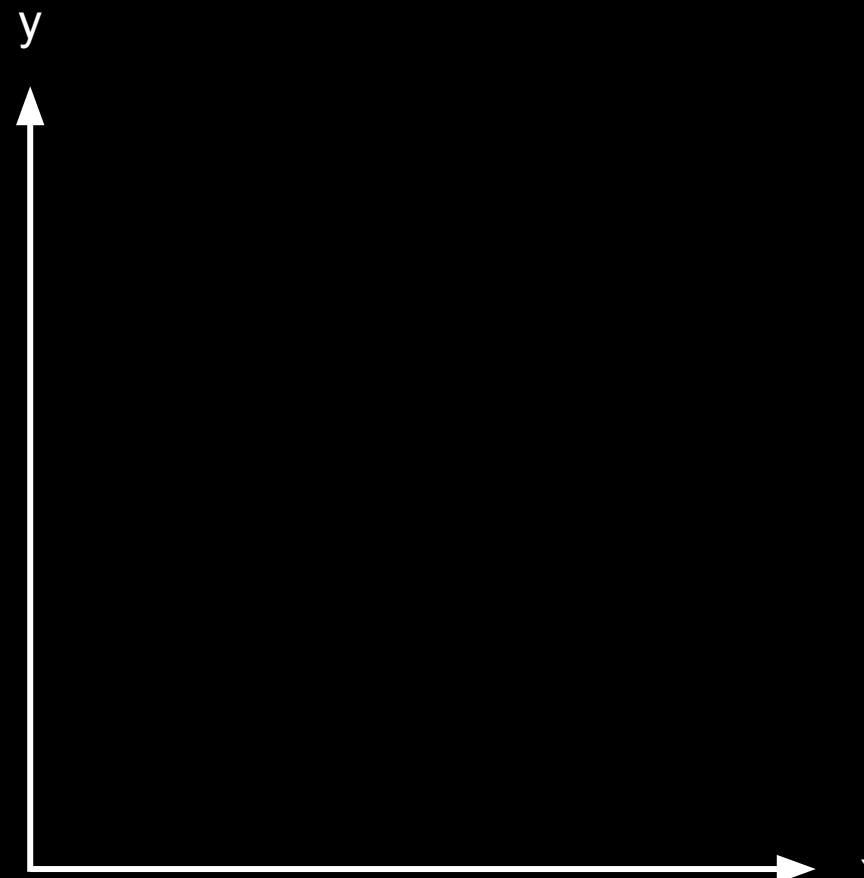


VISUALIZE DATA

data

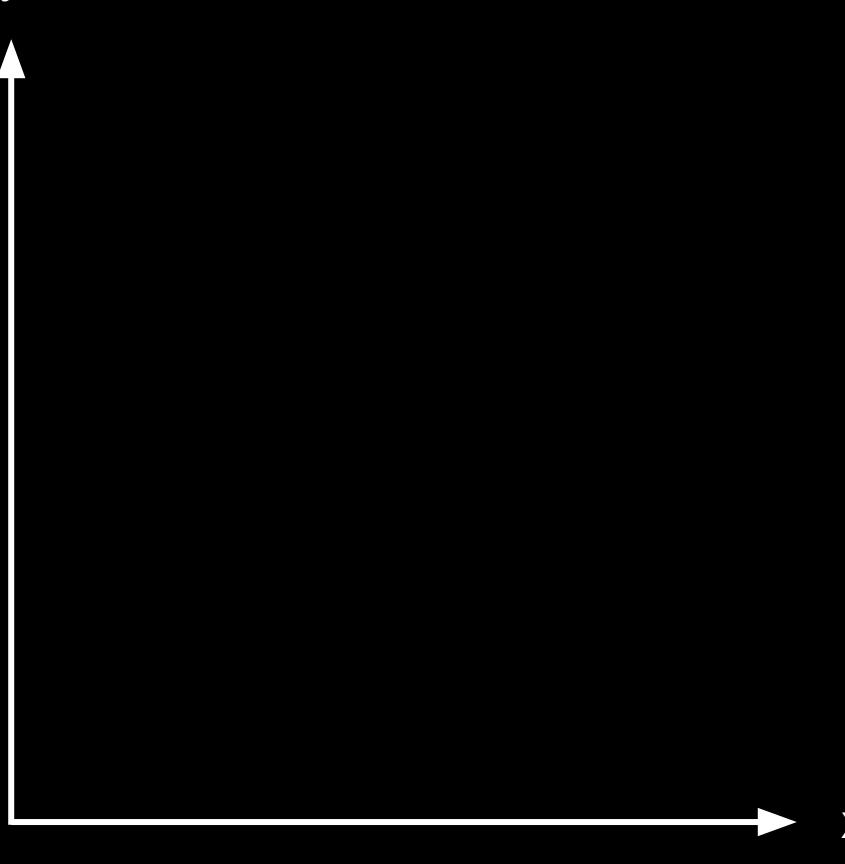
category	pct
A	75
B	33
C	100





y

x



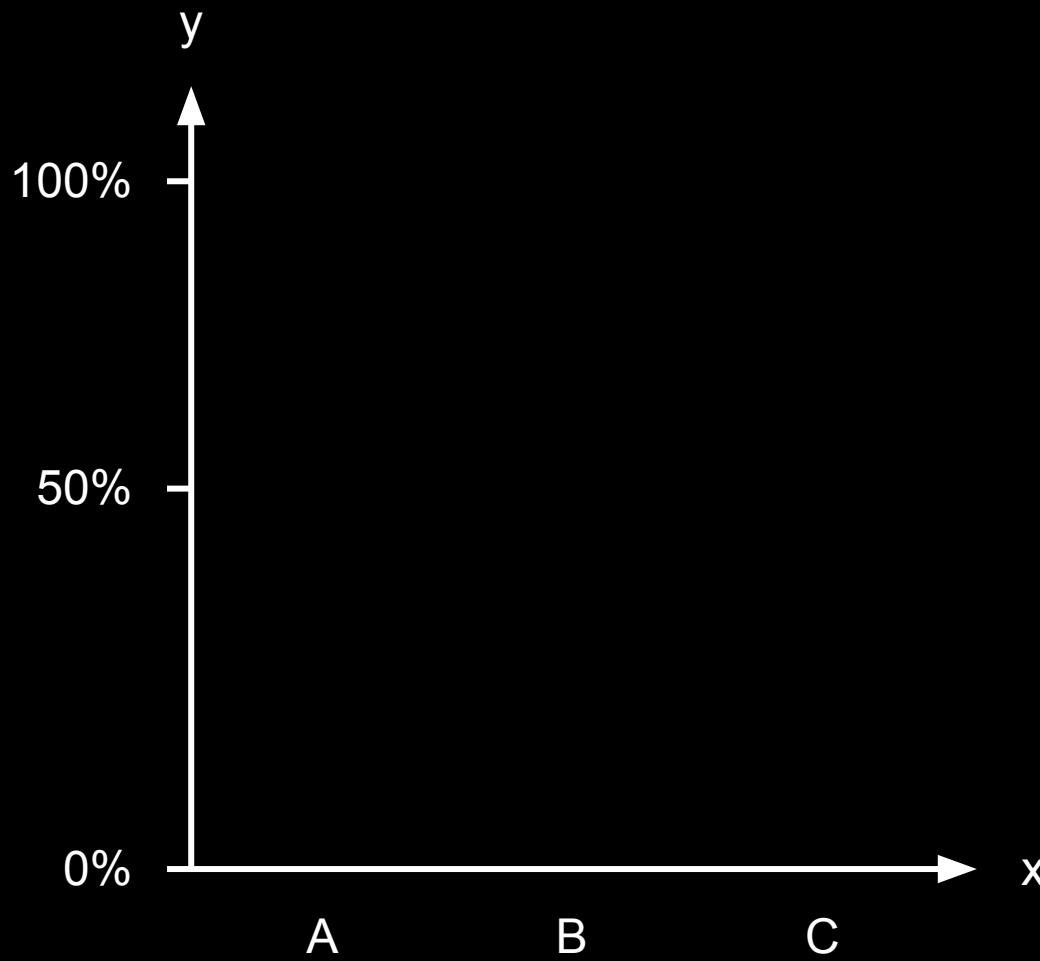
y

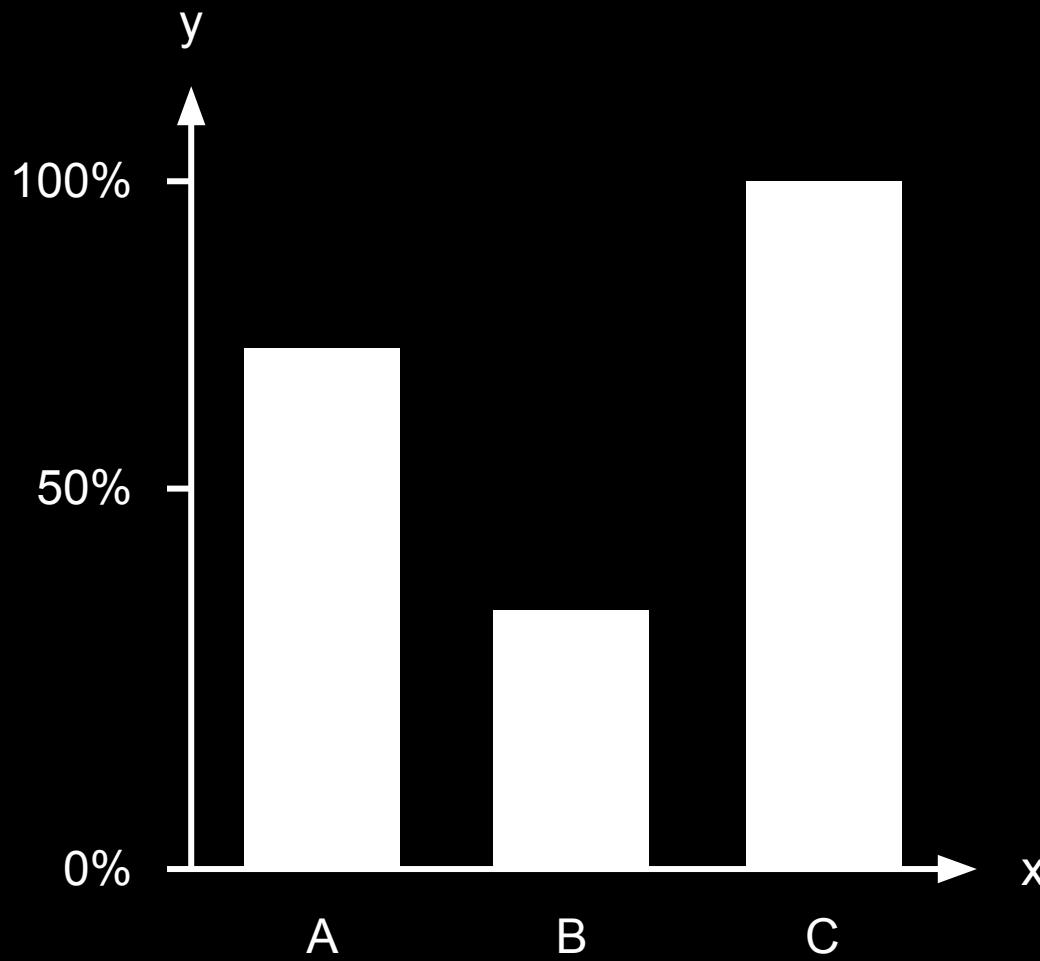
x

A

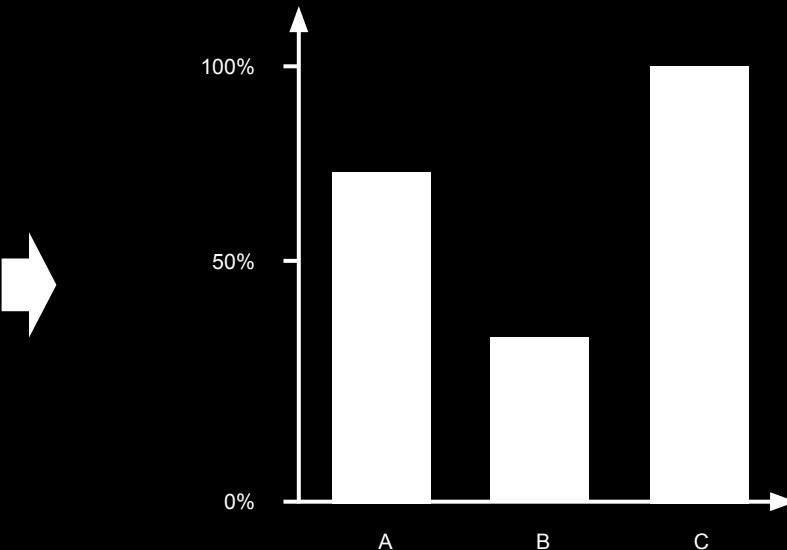
B

C





category	pct
A	75
B	33
C	100



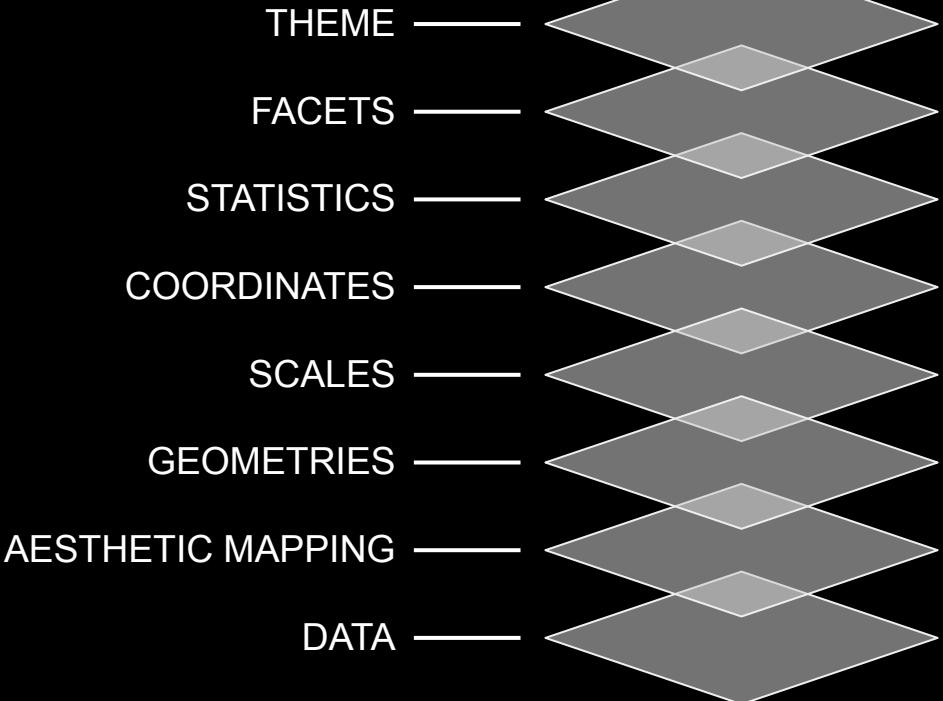
`{{ ggplot2 }}`

why visualize?

`{{ ggplot2 }}`

grammar of graphics

any
data
visualization



has useful defaults



mandatory



THEME

FACETS

STATISTICS

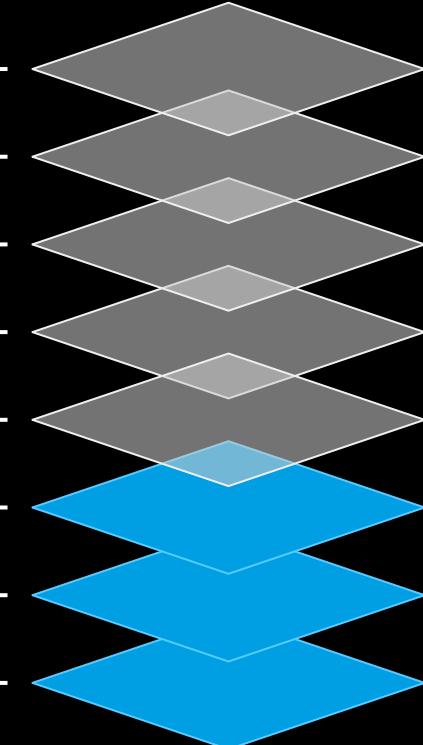
COORDINATES

SCALES

GEOMETRIES

AESTHETIC MAPPING

DATA

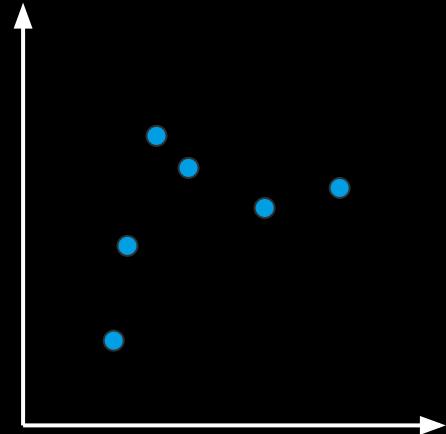


ggplot()

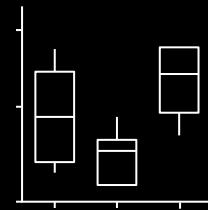
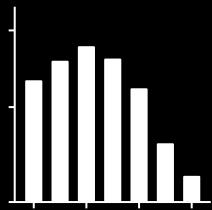
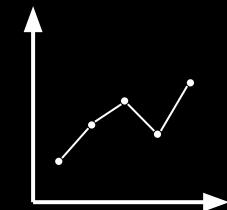
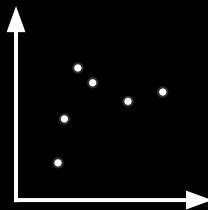
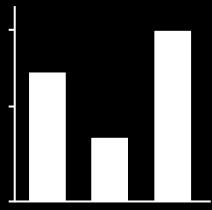
```
ggplot() +  
aes()
```

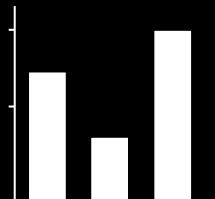
```
ggplot() +  
aes() +  
geom_point()
```

```
ggplot() +  
  aes() +  
  geom_point()
```

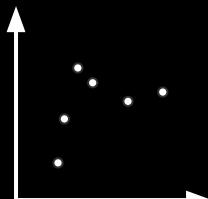


basic plots

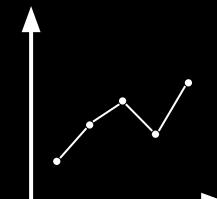




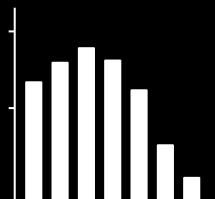
bar chart



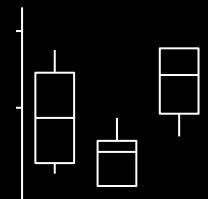
scatter plot



line chart

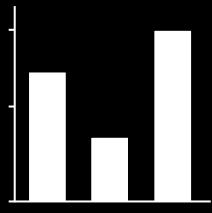


histogram



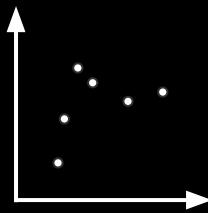
box plot

amounts
proportions
distributions (discrete)



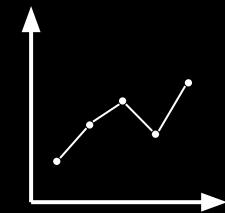
bar chart

associations
patterns



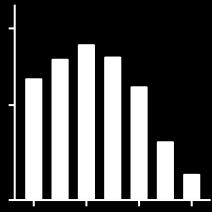
scatter plot

trends
developments



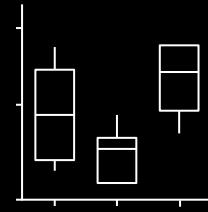
line chart

distributions (continuous)



histogram

compare distributions (continuous)



box plot

COMMUNICATE DATA

quarto

`{{ quarto }}`

markdown

+

R

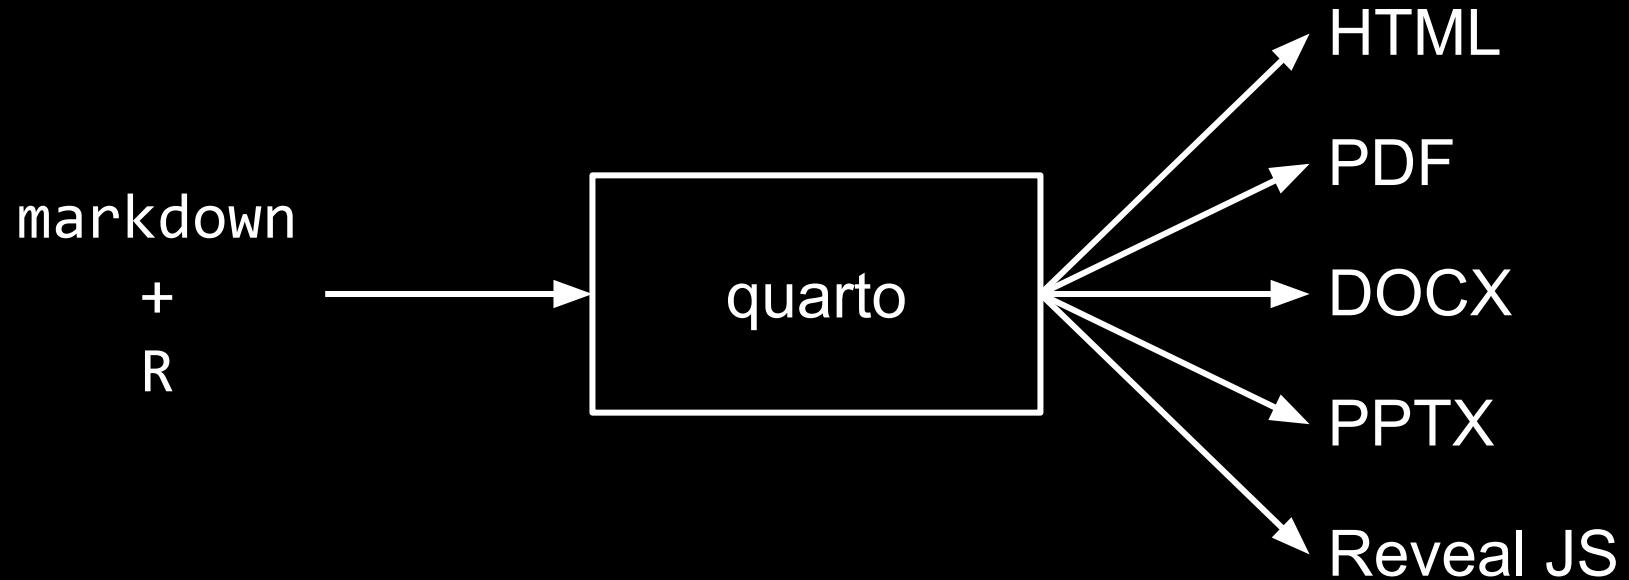
markdown

+

R



quarto



markdown

Heading 1

Heading 2

Heading 3

Heading 4

This is **italic**,
and this ****bold****

This is rendered
as `code`.

- List item A
- List item B
- List item C

1. First
2. Second
3. Third

![Image title](path/to/image.png)

```
![Image title](path/to/image.png){width=200}
```

```
![Image title](path/to/image.png){#fig-myimage}
```

For more details see `@fig-myimage`.

[Linked text](<https://quarto.org>)

`` ` {r}

1 + 1

```

# code options

```
```{r}
```

```
#| echo: false
```

```
1 + 1
```

```
...
```

```
```{r}
```

```
#| eval: false
```

```
x <- 1 + 1
```

```
```
```

```
```{r}
#| include: false
library(tidyverse)
```

```

```
```{r}
#| message: false
data <- read_csv("data.csv")
```

```

```
```{r}
#| warning: false
data <- read_csv("data.csv")
```

```

figures

```
```{r}
#| label: fig-tweets-per-user
#| fig-cap: "Tweets per User"
tweets |>

ggplot() +
 aes(x = screen_name) +
 geom_bar()
```

```

```
```{r}
#| code-fold: true
#| code-summary: "Show code"

tweets |>

 ggplot() +
 aes(x = screen_name) +
 geom_bar()

```
```

cross references

```
```{r}  
#| label: fig-tweets-per-user
#| fig-cap: "Tweets per User"

...
```
```

In `@fig-tweets-per-user` you can see an overview of the number of tweets per user in the data set.

```
# Introduction {#sec-introduction}
```

```
...
```

```
# Analysis
```

As stated in `@sec-introduction`, the goal of this paper is to analyze the user behavior with regard to the content they tweet.

citation & bibliography

<https://quarto.org/docs/authoring/citations.html>

output formats

format: html

format:

html: default

pdf: default

```
format:  
  html:  
    code-fold: true  
  pdf: default
```