



ML-BASED NLP

with spaCy & Python

- The NLP Pipeline
- Getting Started
- NLP Tasks
 - Tokenization
 - Part-of-Speech Tagging
 - Lemmatization
 - Named Entity Recognition
 - Dependency Parsing
 - Stop Words
 - Similarity
 - Sentiment
- Pretrained Language Models

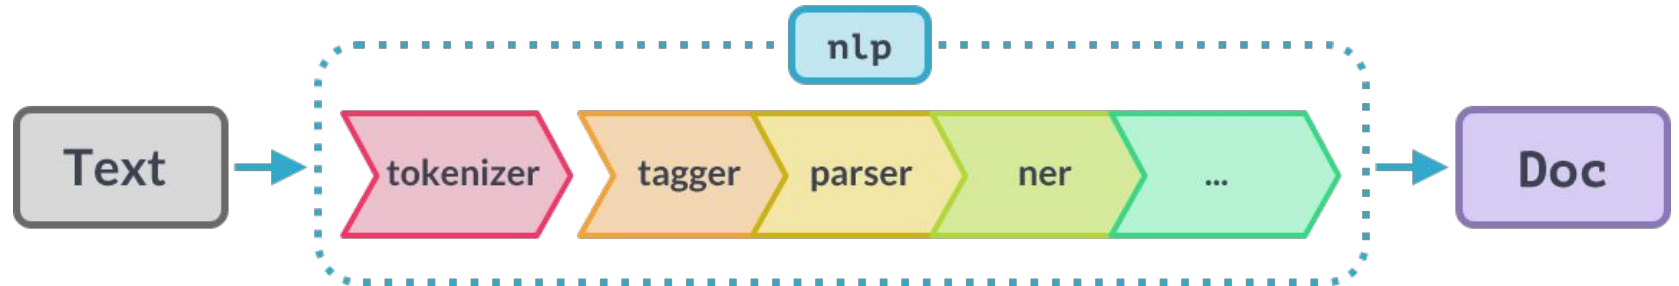
RECOMMENDED RESOURCES

- Free online course “spaCy 101”: <https://spacy.io/usage/spacy-101>
- spaCy’s model documentation website: <https://spacy.io/models>

THE NLP PIPELINE

THE NLP PIPELINE

FROM TEXT TO DOC



Source: <https://spacy.io/usage/spacy-101>

GETTING STARTED

GETTING STARTED

Prerequisite: You need Python 3.x installed on your computer

Step 1: Install spaCy:

```
pip install spacy
```

Step 2: Download a trained model:

```
python -m spacy download en_core_web_sm
```

Step 3: Load spaCy and model from Python:

```
import spacy
```

```
nlp = spacy.load("en_core_web_sm")
```

This is a specific trained model spaCy offers

```
doc = nlp("I love studying at the University of Applied Sciences in Osnabrück")
```

TOKENIZATION


TOKENIZATION

The first component in a NLP pipeline is the tokenizer. It splits the raw text into tokens:

```
doc = nlp("I love studying at the University of Applied Sciences in Osnabrück")
```

```
for token in doc:
```

```
    print(token.text)
```



```
I  
love  
studying  
at  
the  
University  
of  
Applied  
Sciences  
in  
Osnabrück
```

NOTE: The tokenizer is rule-based, and no machine learning is involved here.


PART-OF-SPEECH TAGGING

PART-OF-SPEECH TAGGING

The tagger determines the type of token and its role in the sentence:

```
doc = nlp("I love studying at the University of Applied Sciences in Osnabrück")
```

```
for token in doc:  
    print(token.text, token.pos_, token.tag_)
```



I	PRON	PRP
love	VERB	VBP
studying	VERB	VBG
at	ADP	IN
the	DET	DT
University	PROPN	NNP
of	ADP	IN
Applied	PROPN	NNP
Sciences	PROPN	NNPS
in	ADP	IN
Osnabrück	PROPN	NNP

NOTE: `spacy.explain("NNPS")` will print an explanation → "noun, proper plural"

LEMMATIZER

LEMMATIZER


EXAMPLE 1/2

The lemmatizer determines the base or canonical form (lemma) of a word:

```
doc = nlp("I love studying at the University of Applied Sciences in Osnabrück")
```

```
for token in doc:
```

```
    print(token.text, token.lemma_)
```



```
I I  
love love  
studying study  
at at  
the the  
University University  
of of  
Applied Applied  
Sciences Sciences  
in in  
Osnabrück Osnabrück
```

LEMMATIZER


EXAMPLE 2/2

Given a different context, the word “applied” is converted to its lemma form:

```
doc = nlp("I love data analytics when applied in practice")
```

```
for token in doc:
```

```
    print(token.text, token.lemma_)
```



I	I
love	love
data	data
analytics	analytic
when	when
applied	apply
in	in
practice	practice

NAMED ENTITY RECOGNITION (NER)

NAMED ENTITY RECOGNITION

The NER-component recognizes named entities such as famous people, places, or organisation:

```
doc = nlp("I love studying at the University of Applied Sciences in Osnabrück")
```

```
for entity in doc.ents:  
    print(entity.text, entity.label_)
```



the University of Applied Sciences	ORG
Osnabrück	GPE

DEPENDENCY PARSING


DEPENDENCY PARSING

The dependency parser determines the syntactic relationship between tokens:

```
doc = nlp("I love studying at the University of Applied Sciences in Osnabrück")
```

```
for token in doc:
```

```
    print(token.text, token.dep_, token.head, token.is_sent_start)
```



I	nsubj	love	True
love	ROOT	love	False
studying	xcomp	love	False
at	prep	studying	False
the	det	University	False
University	pobj	at	False
of	prep	University	False
Applied	compound	Sciences	False
Sciences	pobj	of	False
in	prep	University	False
Osnabrück	pobj	in	False

STOP WORDS


STOP WORDS

The NLP-pipeline determines stop words along with other useful characteristics:

```
doc = nlp("I love studying at the University of Applied Sciences in Osnabrück")
```

```
for token in doc:
```

```
    print(token.text, token.is_stop, token.is_alpha, token.shape)
```



I	True	True	X
love	False	True	xxxx
studying	False	True	xxxx
at	True	True	xx
the	True	True	xxx
University	False	True	Xxxxx
of	True	True	xx
Applied	False	True	Xxxxx
Sciences	False	True	Xxxxx
in	True	True	xx
Osnabrück	False	True	Xxxxx

SIMILARITY

SIMILARITY

Models that ship with word vectors (**md**, **lg**) can [calculate similarity scores](#) between tokens and documents:

```
nlp = spacy.load("en_core_web_md")  
doc1 = nlp("I love this course")  
doc2 = nlp("I like this lecture")
```

```
print(doc1, "<->", doc2, doc1.similarity(doc2))
```

I love this course <-> I like this lecture [0.9248678087461984](#)

NOTE: There is no objective definition of “similar” and the meaning depends on the use case.

SENTIMENT
COMING SOON

PRETRAINED LANGUAGE MODELS

- spaCy offers [variety of models](#) for many languages.
- Models differ in **size** and **prediction speed** and **accuracy**:
 - Larger model → slower but higher accuracy
 - Smaller model → faster but less accuracy
- Englisch:
 - **en_core_web_sm** (12 MB) → smallest English model based trained on web-data
 - **en_core_web_md** (40 MB) → medium English model, contains word vectors
 - **en_core_web_lg** (560 MB) → large English model, highest accuracy
 - **en_core_web_trf** (438 MB) → Transformer-based model, GPU recommended